

# CS 344: OPERATING SYSTEMS I

## 01.23: THREADS

M/W 12:00 – 1:50 PM (LINC #200)

Sanghyun Hong

[sanghyun.hong@oregonstate.edu](mailto:sanghyun.hong@oregonstate.edu)



**Oregon State**  
University

**SAIL**

Secure AI Systems Lab

# NOTICE

---

- Deadlines
  - (1/23 11:59 PM) Programming assignment 1
  - (1/30 11:59 PM) Midterm quiz 1
  - (2/06 11:59 PM) Programming assignment 2

# RECAP

---

- Part I: Process
  - Provide abstraction
    - What is a program?
    - What is a process?
    - How does OS run a program?
  - Offer standard libraries
    - How do we run (or stop) a process?
    - How does OS manage the process(es) we ran?
  - Manage resources
    - (Note) We will talk about this in the “scheduling” class

# PRACTICE QUESTION: C

- **Static variables, functions, etc.**

- What will be the “google” stock price?
- What will be the prices of both stocks?
- What will be printed out to Terminal?
- What will be the prices of both stocks?
- What will be printed out to Terminal?

Required headers...

```
static int google_stock = 2000;
```

```
int increase_price(int stock_price, int amount) {  
    stock_price = stock_price + amount;  
    google_stock = google_stock + amount;  
}
```

```
int main(void) {
```

Run --> 

```
    int apple_stock = 99;
```

```
    printf("Google stock price is %d\n", google_stock);
```

Run --> 

```
    printf("Apple stock price is %d\n", apple_stock++);
```

```
    apple_stock = increase_price(apple_stock, 50);
```

```
    printf("Google stock price is %d\n", google_stock);
```

Run --> 

```
    printf("Apple stock price is %d\n", ++apple_stock);
```

```
    return 0;
```

```
}
```

# PRACTICE QUESTION: C

---

- **Pointers and strings**

- What will be the value of “ilen”?
- What will be the value of “slen”?
- How many bytes “str” uses in memory?
- What will be the execution result?

Required headers...

```
int main(void) {  
    int slen = 0;  
    int *iptr = NULL;  
    char str = "Hello world!";  
  
    Run -->    ilen = sizeof(iptr);  
    Run -->    slen = strlen(str);  
    Run -->    printf("The length of this string is %d\n", slen);  
  
    return 0;  
}
```

# PRACTICE QUESTION: C

- **Pointer operations**

- What will be printed out to Terminal?
- What will be printed out to Terminal?
- What will be printed out to Terminal?
- What will be printed out to Terminal?

Required headers...

```
int swap(int num1, int *num2) {  
    int temp = num1;  
    num1 = num2;  
    *num2 = temp;  
    return num1;  
}
```

```
int main(void) {  
    int val1 = 1;  
    int val2 = 2;  
    int vals[] = { 10, 20, 30, 40, 50 };  
    int *ptr = vals;
```

Run --> printf("Val1 / 2 / 3: %d, %d, %d\n", val1, val2, vals[0]);

Run --> printf("Val1 / 2: %d, %d\n", \*(ptr+2), (\*ptr)+2);

```
    val1 = swap(val1, ptr);
```

Run --> printf("Val1 / 2 / 3: %d, %d, %d\n", val1, val2, vals[0]);

Run --> printf("Val1 / 2: %d, %d\n", \*(ptr+2), (\*ptr)+2);

```
    return 0;
```

```
}
```

# PRACTICE QUESTION: PROCESS

- **Segments (components) of a process**

- Which segment “counter1” is?
- Which segment “ret” is?
- Which segment “counter2” is?
- Which segment “buf” is?
- What are the counter1 and 2 values?
- Which segment “ret” is?
- Which segment “buf” is?

Required headers...

```
#define BUFSIZE 512
```

```
static int counter1 = 0;
```

```
int my_function() {
```

```
Run --> int counter2 = 2;
```

```
Run --> char *buf = (char *) malloc(BUFSIZE * sizeof(char));
```

```
counter1 = counter1 + 1;
```

```
counter2 = counter2 - 1;
```

```
Run --> return counter2;
```

```
}
```

```
int main(void) {
```

```
Run --> int ret = 0;
```

```
Run --> ret = my_function();  
printf("Ret: %d\n", ret);
```

```
return 0;
```

```
}
```

# TOPICS FOR TODAY

---

- Part I: Threads
  - Provide abstraction
    - What is a thread?
    - How is it different from a process?
    - How does OS run threads?
  - Offer standard libraries
    - How do we create/run/kill a thread?
    - How does OS manage the thread(s) we ran?
  - Manage resources
    - (Note) We will talk about this in the “scheduling” and “synchronization” classes



# RUNNING MULTIPLE PROCESSES: WEB-SERVER EXAMPLE

---

- Amazon.com:
  - What does the webserver do?

# WEB-SERVER EXAMPLE

- Amazon.com:
  - A user requests the website
  - A server accepts the connection
  - A server sends the webpage to the user
  - A user clicks something
  - A server sends the webpage as a response
  - ... (continue)

## Pseudo code (server)

```
int main(void) {  
    // 1. server accepts the connection  
    connection = accepts(user-request, ...)  
  
    // 2. server sends the webpage to user  
    sends_webpage(connection, html-page)  
  
    // 3. server starts accepting the user requests  
    while (action = receive_request(connection)) {  
        if (action == login) {  
            if (!correct_credential(action.id, action.pw))  
                return -1; // return error, login fail  
            connection.login_success = 1;  
        }  
  
        ....  
    }  
  
    return 0; // halt the webserver, never reached  
}
```

# WEB-SERVER EXAMPLE – CONT'D

- Amazon.com:
  - A user requests the website
  - A server accepts the connection
  - A server sends the webpage to the user
  - A user clicks something
  - A server sends the webpage as a response
  - ... (continue)

**What would be a potential problem?**

## Pseudo code (server)

```
int main(void) {  
    // 1. server accepts the connection  
    connection = accepts(user-request, ...)  
  
    // 2. server sends the webpage to user  
    sends_webpage(connection, html-page)  
  
    // 3. server starts accepting the user requests  
    while (action = receive_request(connection)) {  
        if (action == login) {  
            if (!correct_credential(action.id, action.pw))  
                return -1; // return error, login fail  
            connection.login_success = 1;  
        }  
  
        ....  
    }  
  
    return 0; // halt the webserver, never reached  
}
```

# WEB-SERVER EXAMPLE – CONT'D

- Amazon.com:

- A user requests the website
- A server accepts the connection
- A server sends the webpage to the user
- A user clicks something
- A server sends the webpage as a response
- ... (continue)

→ This procedure will be the **same** for all users  
> **Multi-process** web-server

## Pseudo code (server)

```
int main(void) {  
    // 1. server accepts the connection  
    connection = accepts(user-request, ...)  
  
    // 2. server sends the webpage to user  
    sends_webpage(connection, html-page)  
  
    // 3. server starts accepting the user requests  
    while (action = receive_request(connection)) {  
        if (action == login) {  
            if (!correct_credential(action.id, action.pw))  
                return -1; // return error, login fail  
            connection.login_success = 1;  
        }  
  
        ....  
    }  
  
    return 0; // halt the webserver, never reached  
}
```

# MULTI-PROCESS WEB-SERVER EXAMPLE

- Amazon.com:

- A user requests the website
- A server accepts the connection
- A server sends the webpage to the user
- A user clicks something
- A server sends the webpage as a response
- ... (continue)


This procedure will be the **same** for all users  
> **Multi-process** web-server

## Pseudo code (server)

```
int main(void) {  
    while(connection = accepts(user-request, ...)) {  
        // fork: create a new process  
        switch(pid = fork()) {  
            case 0:  
                // server sends the webpage to user  
                sends_webpage(connection, html-page)  
  
                // server starts accepting the user requests  
                while (action = receive_request(connection)) {  
                    if (action == login) {  
                        if (!correct_credential(action.id, action.pw))  
                            return -1; // return error, login fail  
                        connection.login_success = 1;  
                    }  
                    ....  
                }  
            }  
        } // end of switch ...  
    } ...  
}
```

# MULTI-PROCESS WEB-SERVER EXAMPLE: OS VIEW

- Amazon.com:
  - A user requests the website
  - A server accepts the connection
  - A server sends the webpage to the user
  - A user clicks something
  - A server sends the webpage as a response
  - ... (continue)

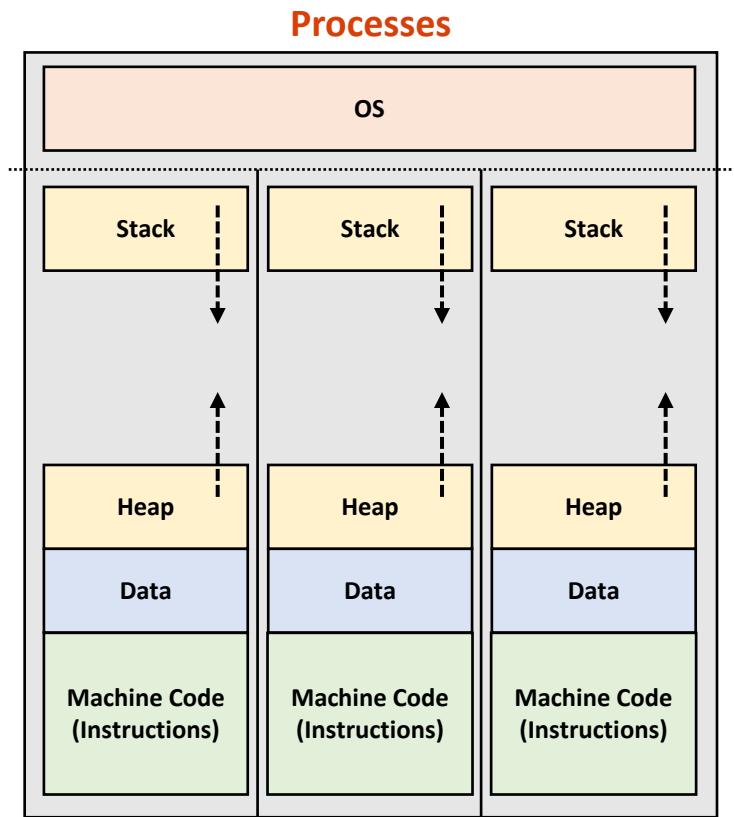


# MULTI-PROCESS WEB-SERVER EXAMPLE: POTENTIAL ISSUES

- Data is not shared between processes
  - A user requests the website
  - ... (continue)

```
int main(void) {  
  
    // initialize some data in this block  
  
    while(connection = accepts(user-request, ...)) {  
        // fork: create a new process  
        switch(pid = fork()) {  
            case 0:  
                // server sends the webpage to user  
                sends_webpage(connection, html-page)  
  
                ...  
        }  
    }  
}
```

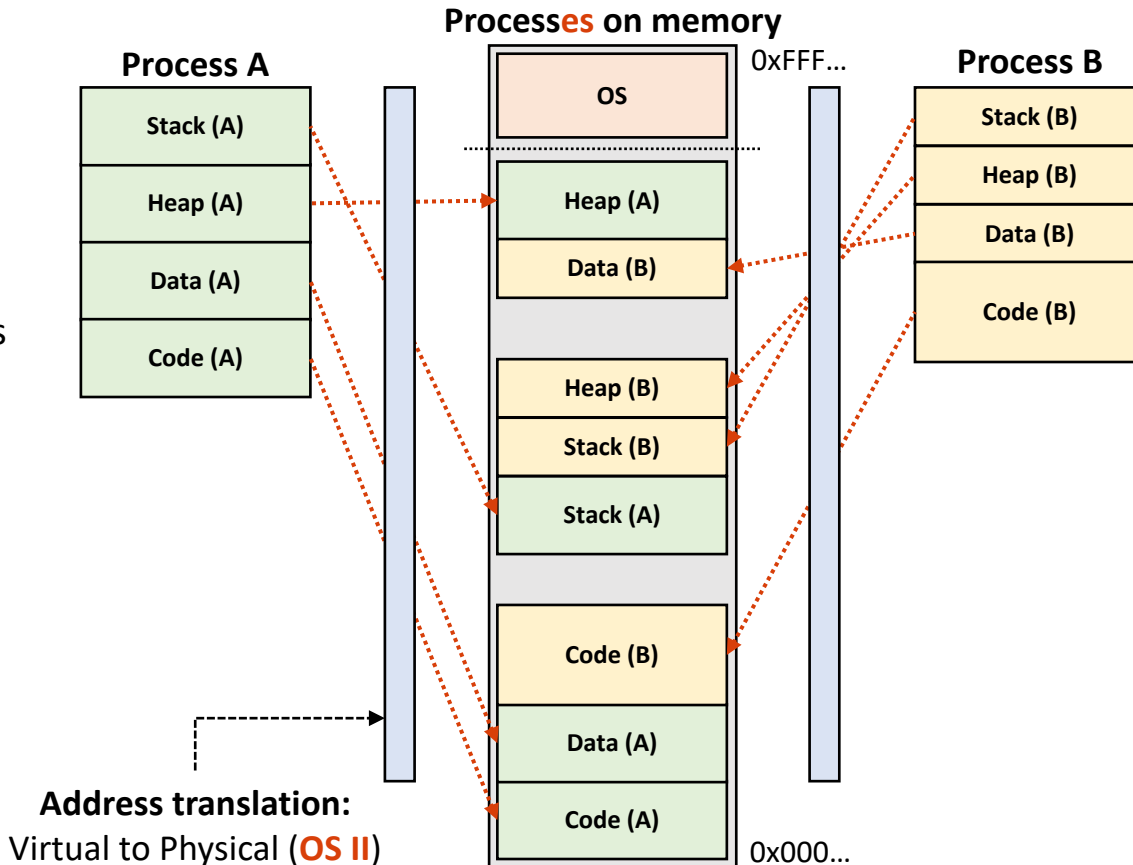
The data in the above block **won't be shared** between processes; each process will have a copy of the same data (\*causes memory overhead)



# NOTE: WHY ISN'T THE DATA SHARED BETWEEN PROCESSES?

- **Process isolation**

- No segment is shared
- Security reasons
  - Data breach
  - System crashes
  - Control other processes
  - ...
- **Access: seg-faults!**





# **SOLUTION: THREADS**

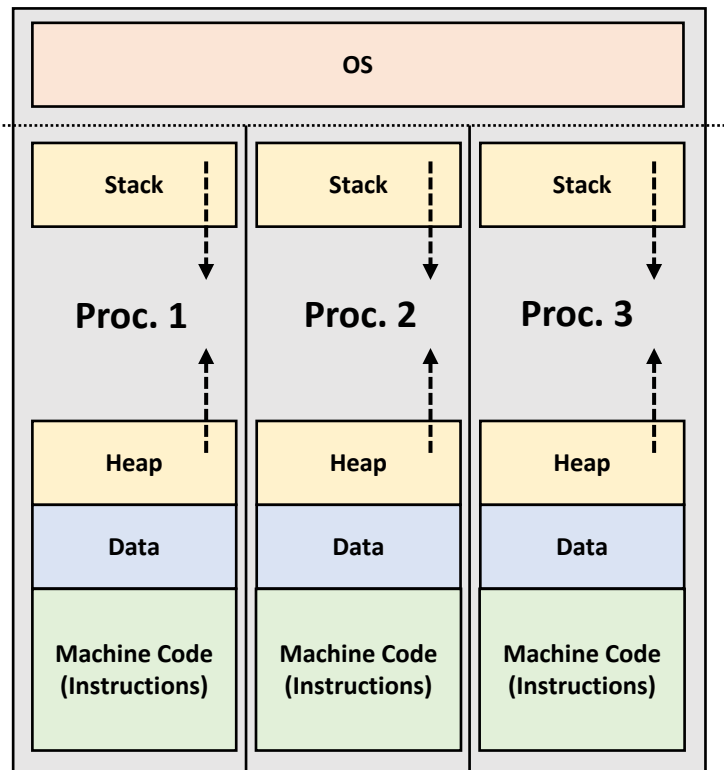
# PROVIDE ABSTRACTION: A THREAD

---

- Thread
  - **Definition:** a smallest schedulable execution context
  - **Terminology:**
    - Smallest: it's much light-weight than a process
    - Schedulable execution context: one thread can run on a CPU at a time

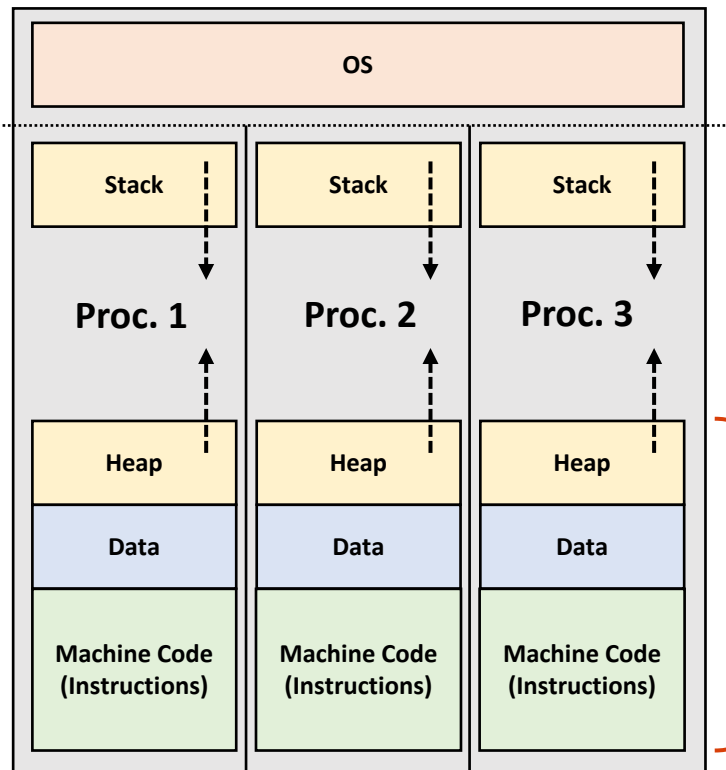
# PROVIDE ABSTRACTION: A THREAD – CONT'D

Processes on memory

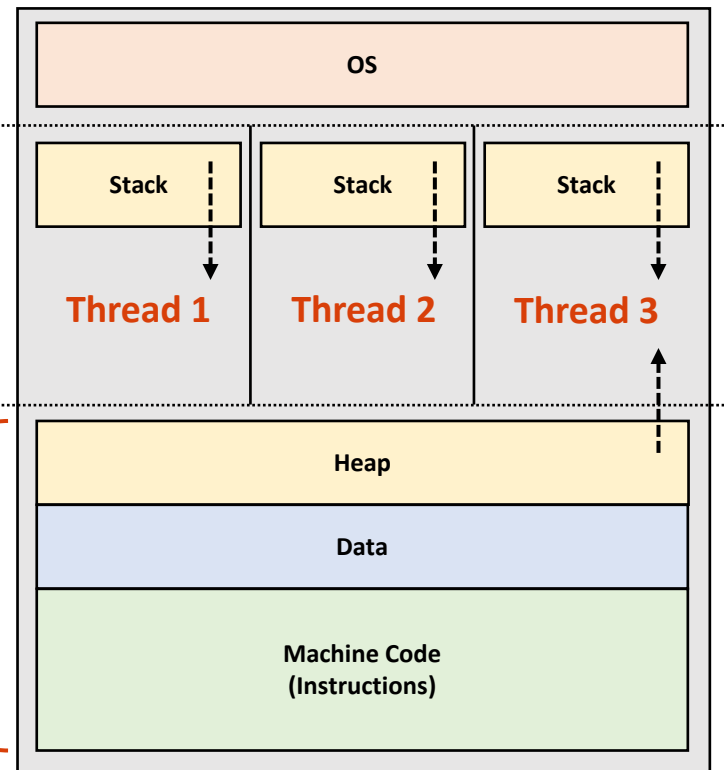


# PROVIDE ABSTRACTION: A THREAD – CONT'D

**Processes** on memory



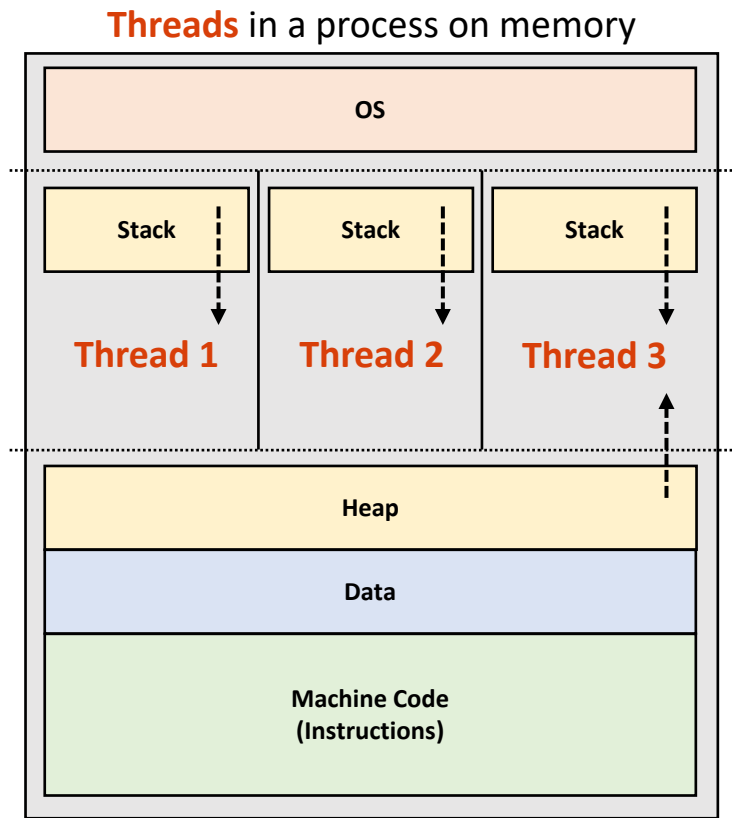
**Threads** in a process on memory



**Reduce  
Duplications**

# PROVIDE ABSTRACTION: HOW IS IT DIFFERENT FROM A PROCESS?

- Threads share:
  - **Code** and **data** segments
  - **Heap** memory (ex. global variables)
  - Open files (ex. I/O access points)
- Threads **do not** share:
  - **Stack** segments, e.g.:
    - arguments passed when we launch them
    - local variables we initialize within them
    - return address, when they terminate ([OS II](#))
  - Running contexts, e.g.:
    - process state
    - stack pointer
    - ...



# PROVIDE ABSTRACTION: HOW OS DEFINES A THREAD?

- (Linux) has “thread control block”

- Code

- Program counter
    - Instruction pointer

- Stack and heap

- Stack pointer
    - Heap pointer

- Running context

- Process state (ID, ...)
    - Execution flags
    - CPU # to run
    - (OS II) Scheduling policy
    - (OS II) Mem. virtualization

**Thread Control Block:** A set of information that OS requires to run a thread on a CPU, different from CPU vendors (ex. In Linux, it's the same: *task\_struct*, [Link](#))

```
... 728 struct task_struct {
729     #ifdef CONFIG_THREAD_INFO_IN_TASK
730         /*
731          * For reasons of header soup (see current_thread_info()), this
732          * must be the first element of task_struct.
733          */
734         struct thread_info      thread_info;
735     #endif
736     unsigned int                __state;
737
738     #ifdef CONFIG_PREEMPT_RT
739         /* saved state for "spinlock sleepers" */
740         unsigned int            saved_state;
741     #endif
742
743     /*
744      * This begins the randomizable portion of task_struct. Only
745      * scheduling-critical items should be added above here.
746      */
747     randomized_struct_fields_start
748
749     void *stack;
750     refcount_t usage;
751     /* Per task flags (PF_*), defined further below: */
752     unsigned int flags;
753     unsigned int ptrace;
754
755     struct sched_info            sched_info;
756
757     struct list_head             tasks;
758     #ifdef CONFIG_SMP
759     struct plist_node             pushable_tasks;
760     struct rb_node                pushable_dl_tasks;
761     #endif
762
763     struct mm_struct              *mm;
764     struct mm_struct              *active_mm;
765
766     /* Per-thread vma caching: */
767     struct vmacache                vmacache;
768
769     #ifdef SPLIT_RSS_COUNTING
770     struct task_rss_stat            rss_stat;
771     #endif
772
773     int exit_state;
774     int exit_code;
775     int exit_signal;
776     /* The signal sent when the parent dies: */
777     int pdeath_signal;
778     /* JOBCTL_*, siglock protected: */
779     unsigned long jobctl;
780
781     /* Previous Linux versions: */
782     ...
783 }
```

**A process and a thread are the same for OS**

# TOPICS FOR TODAY

---

- Part I: Threads
  - Provide abstraction
    - What is a thread?
    - How is it different from a process?
    - How does OS run threads?
  - Offer standard libraries
    - How do we create/run/kill a thread?
    - How does OS manage the thread(s) we ran?
  - Manage resources
    - (Note) We will talk about this in the “scheduling” and “synchronization” classes

# OFFER STANDARD INTERFACE

---

- How do we run a thread?
  - **System calls**
  - OS provide a set of system calls to control thread execution




# OFFER STANDARD INTERFACE: **THREAD-SPECIFIC** SYSTEM CALLS

---


- Thread-specific system calls
  - **pthread\_create**(thread, attribute, subroutine, subroutine-arguments);
    - Create a new thread executing the *subroutine* in the current process
    - Returns zero if it's successful, otherwise it returns [errno](#)
  - **pthread\_exit**(return-value);
    - Terminate the thread and returns the *return-value* to any successful join
    - Note: If a thread terminates, it will be automatically called and always return success
  - **pthread\_join**(thread, return-value-loc);
    - Suspend execution of the calling thread until the *thread* terminates
    - Once the thread terminates, the function will copy the return value to *return-value-loc*
    - Returns zero if it's successful, otherwise it returns an error

# OFFER STANDARD INTERFACE: **THREAD-SPECIFIC** SYSTEM CALLS

- System call
  - **Example:** `pthread_create(...)`



# THREAD PROGRAMMING: **FORK-JOIN** PATTERN



- **Fork - Join** Pattern
  - **Fork**: Main process creates a set of sub-(or child)-threads that runs a function
  - Each thread exits if the function returns
  - **Join**: Main waits until all the threads exit
- **Example**: download a large file
  - Splits a file into smaller chunks
  - Create a thread for downloading each
  - Sum-up all the downloaded chunks and combine them to create a single large file

# OFFER STANDARD INTERFACE: **THREAD-SPECIFIC** SYSTEM CALLS

- Thread sample code in C
  - How many threads are there?
  - Which thread is created first?
  - Which thread is created last?
  - Which thread runs first/last?
  - What'd be an order of thread joins?
  - What will happen if we run this again?

```
static int value = 128;

void *subroutine (void *threadid) {
    long tid = (long) threadid;
    printf("Thread ID [%lx], value [%d]\n", tid, value ++);
}

int main(int argc, char *argv[]) {
    long t;
    int nthreads = 3;

    pthread_t *threads = (pthread_t *) malloc(nthreads * sizeof(pthread_t));
    memset(threads, 0x00, nthreads * sizeof(pthread_t));

    for (t = 0; t < nthreads; t++) {
        int rc = pthread_create(&threads[t], NULL, subroutine, (void *)t);
        if (rc) {
            printf("[Error] return code is: %d, abort.\n", rc);
            exit(-1);
        }
    }

    for (t = 0; t < nthreads; t++)
        pthread_join(threads[t], NULL);

    return 0;
}
```

# OFFER STANDARD INTERFACE: **THREAD-SPECIFIC** SYSTEM CALLS

- Thread sample code in C
  - How many threads are there?
  - Which thread is created first?
  - Which thread is created last?
  - Which thread runs first/last?
  - What'd be an order of thread joins?
  - What will happen if we run this again?

## Possible execution result:

Thread ID [0], value [128]

Thread ID [2], value [129]

Thread ID [1], value [130]

```
static int value = 128;
```

```
void *subroutine (void *threadid) {  
    long tid = (long) threadid;  
    printf("Thread ID [%lx], value [%d]\n", tid, value++);  
}
```

```
int main(int argc, char *argv[]) {  
    long t;  
    int nthreads = 3;
```

```
    pthread_t *threads = (pthread_t *) malloc(nthreads * sizeof(pthread_t));  
    memset(threads, 0x00, nthreads * sizeof(pthread_t));
```

```
    for (t = 0; t < nthreads; t++) {  
        int rc = pthread_create(&threads[t], NULL, subroutine, (void *)t);  
        if (rc) {  
            printf("[Error] return code is: %d, abort.\n", rc);  
            exit(-1);  
        }  
    }  
}
```

```
    for (t = 0; t < nthreads; t++)  
        pthread_join(threads[t], NULL);
```

```
    return 0;  
}
```


# OFFER STANDARD INTERFACE: HOW OS MANAGES THREADS?

---

- **(Linux) OS**
  - A thread is treated as the same as a process
  - (Linux) thread control block  $\approx$  process context
- A thread can have **three states**:
  - **Ready**: a thread is created and ready to run, but not running now
  - **Running**: a thread running now
  - **Blocked**: a thread is unable to run (terminated or errors)

# OFFER STANDARD INTERFACE: HOW OS MANAGES THREADS?

- Mem layout with two threads
  - Each thread has its own stack
  - Data, code and heap are shared between the two



# TOPICS FOR TODAY

---

- Part I: Threads
  - Provide abstraction
    - What is a thread?
    - How is it different from a process?
    - How does OS run threads?
  - Offer standard libraries
    - How do we create/run/kill a thread?
    - How does OS manage the thread(s) we ran?
  - Manage resources
    - (Note) We will talk about this in the “scheduling” and “synchronization” classes



# Thank You!

M/W 12:00 – 1:50 PM (LINC #200)

Sanghyun Hong

[sanghyun.hong@oregonstate.edu](mailto:sanghyun.hong@oregonstate.edu)



**Oregon State**  
University

**SAIL**

Secure AI Systems Lab