

CS 344: OPERATING SYSTEMS I

02.20: PART III: NETWORKING I

M/W 12:00 – 1:50 PM (LINC #200)

Sanghyun Hong

sanghyun.hong@oregonstate.edu



Oregon State
University

SAIL

Secure AI Systems Lab

NOTICE

- Announcements
 - No lecture on the 27th
 - A slot for quizzes, assignments, and extra opportunities
 - SH will be on Discord
 - 2 more extra credit opportunities on Canvas
 - Build an ML classifier (+2%)
 - Multi-process data loader (+3%)

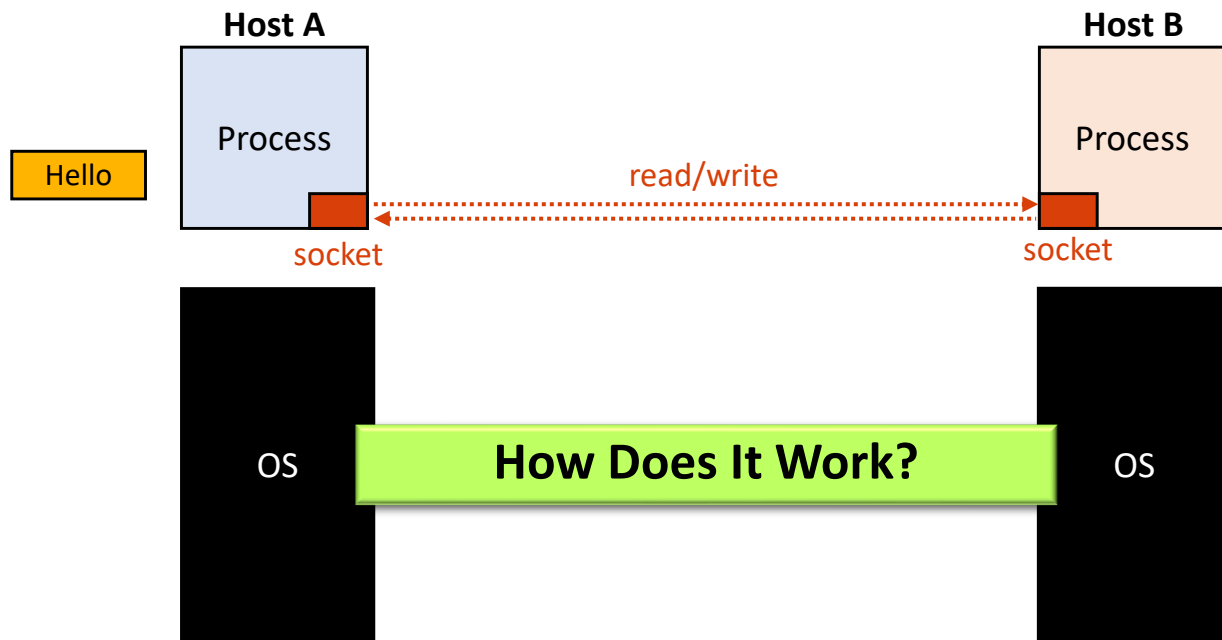
TOPICS FOR TODAY

- Part III: Networking
 - Provide abstraction
 - OSI models
 - Packet encapsulation
 - Offer standard interface
 - RPC mechanisms (e.g., sockets)
 - Manage resources
 - Packet encapsulation in detail

(COMPUTER) NETWORKING

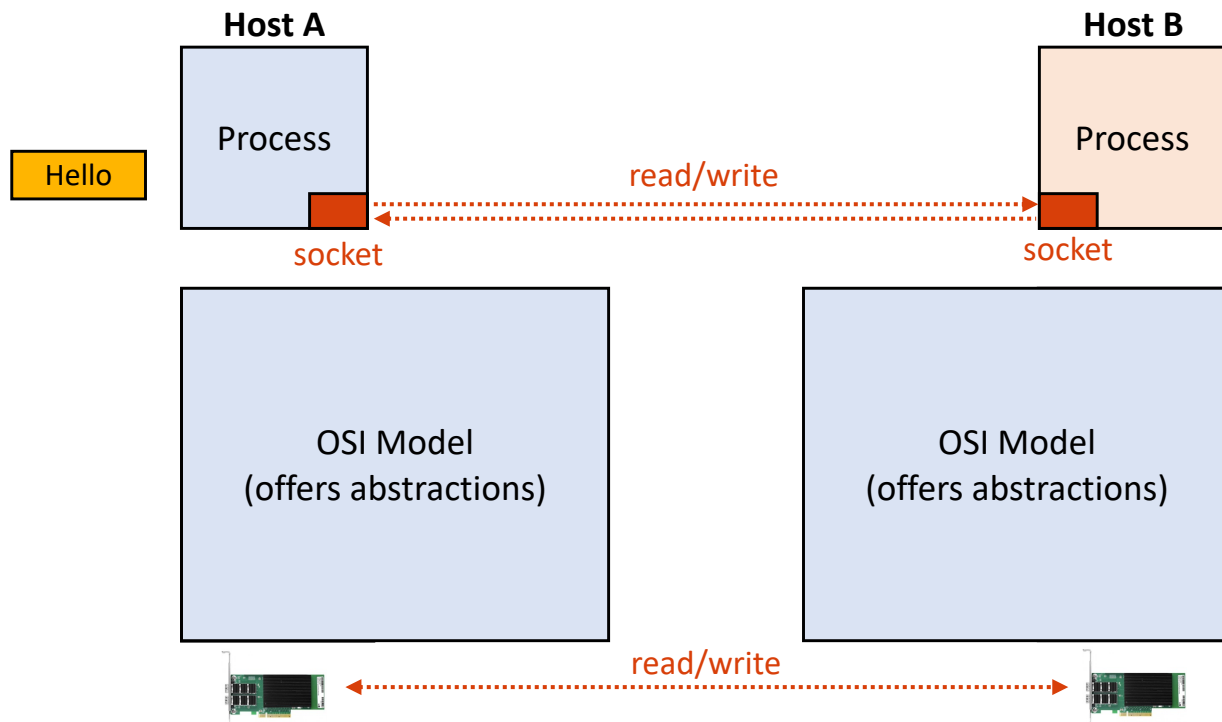
- Networking

- **Definition:** two or more applications on different computers (hosts) exchanging data



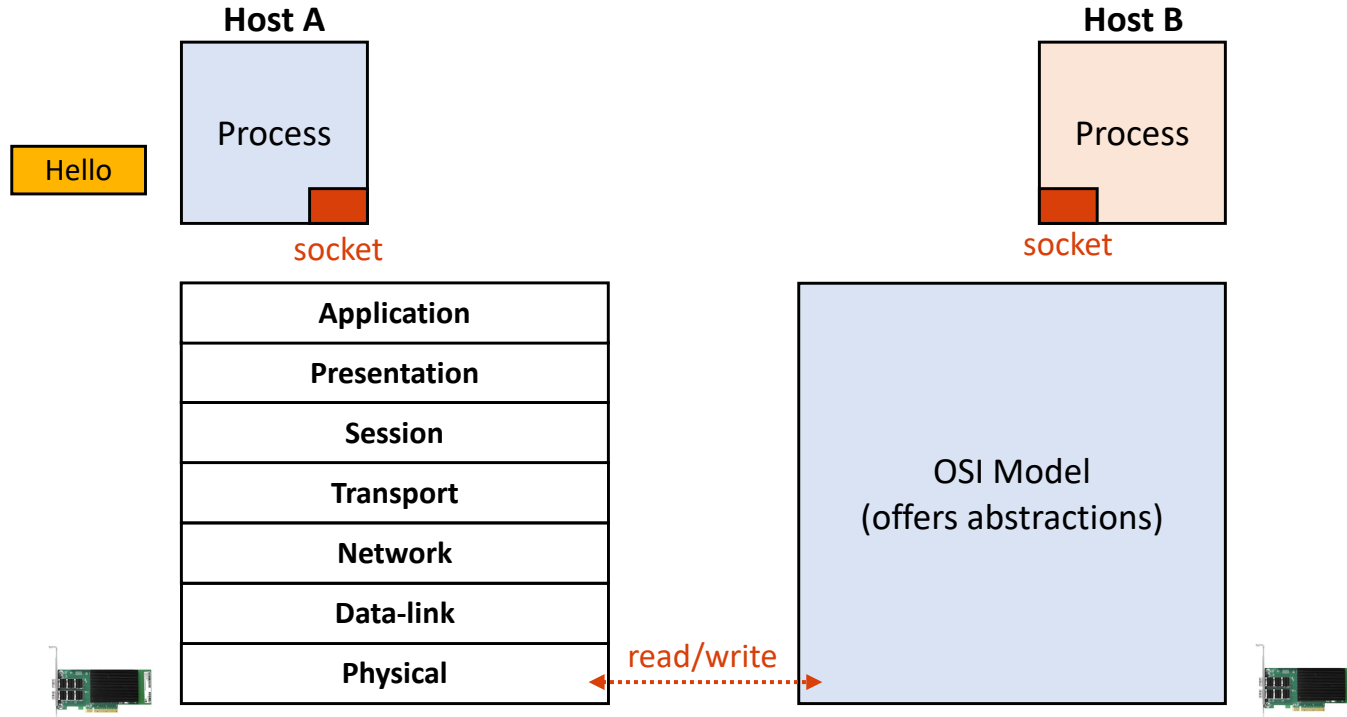
PROVIDE ABSTRACTION

- Open Internet Interface (**OSI**) model



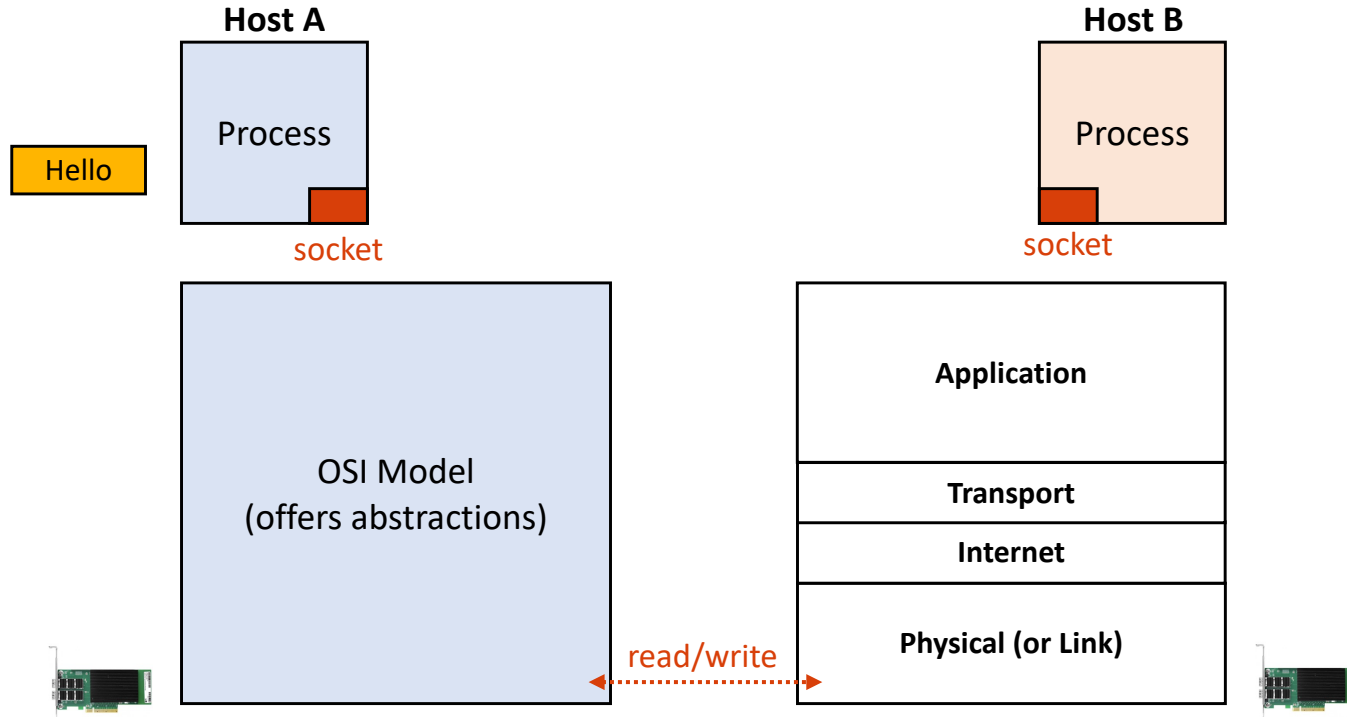
PROVIDE ABSTRACTION: 7-LAYER MODEL

- Open Internet Interface (OSI) model



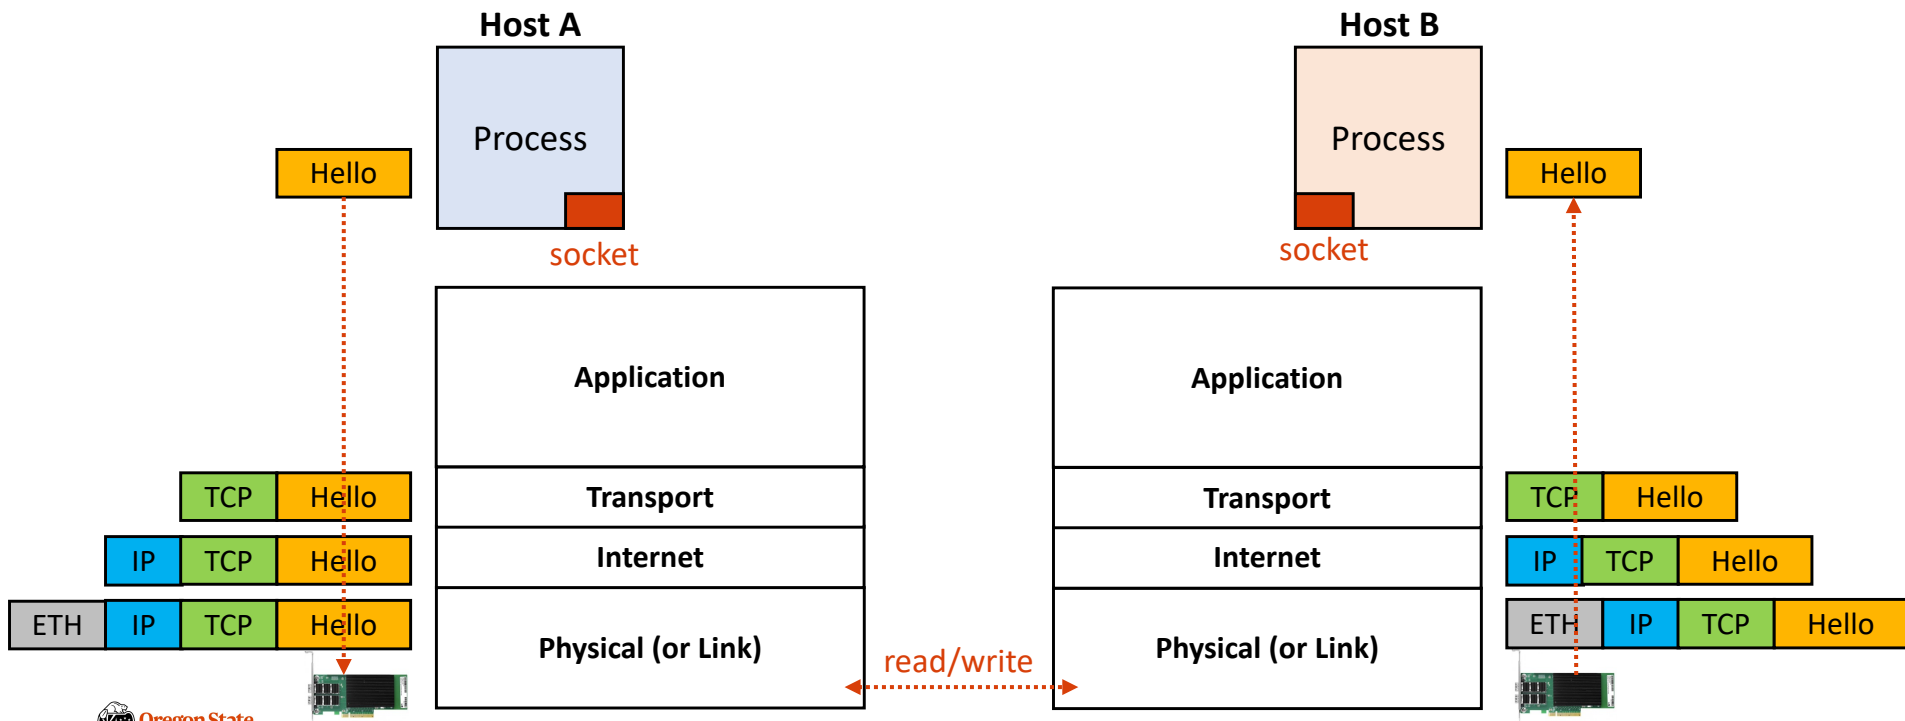
PROVIDE ABSTRACTION: TCP/IP 4-LAYER MODEL

- Open Internet Interface (**OSI**) model



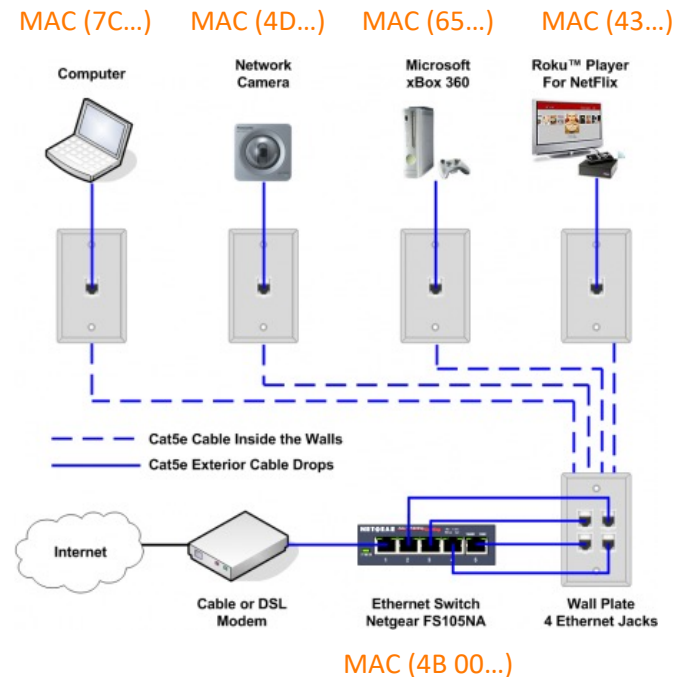
PROVIDE ABSTRACTION: PACKET ENCAPSULATION

- In the TCP/IP 4-layer model



PROVIDE ABSTRACTION: ETHERNET (PHYSICAL LAYER)

- Ethernet Protocol (~80s)
 - Each network device (NIC) has 48-bit **MAC address**
 - Each NIC is connected via Ethernet **cable**

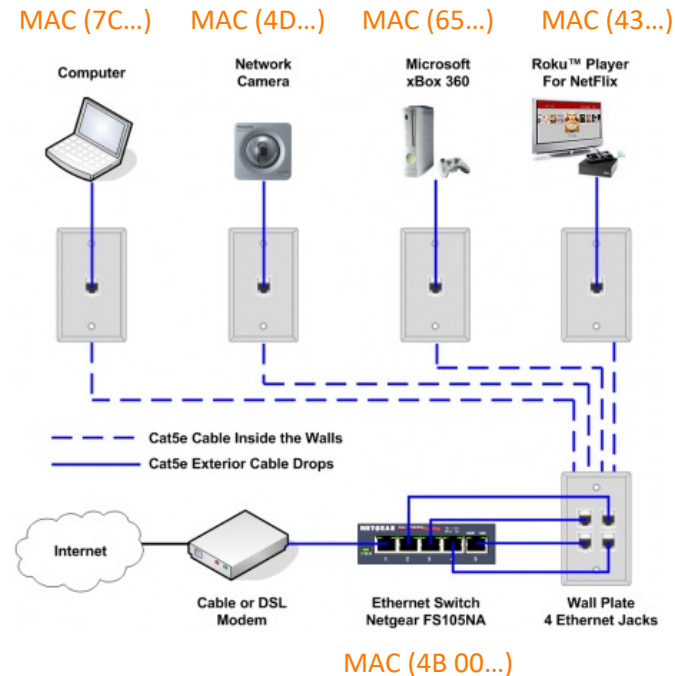


PROVIDE ABSTRACTION: ETHERNET (PHYSICAL LAYER)

- Ethernet Protocol (~80s)

- Each network device (NIC) has 48-bit **MAC address**
- Each NIC is connected via Ethernet **cable**
- **ETH header** contains:

- (64 bit) Preamble (0x11111111... or a unique data)
- (48-bit) Destination MAC address
- (48-bit) Source MAC address
- (16-bit) Type
- (up to 1500 bytes) Data
- (32-bit) CRC for error correcting

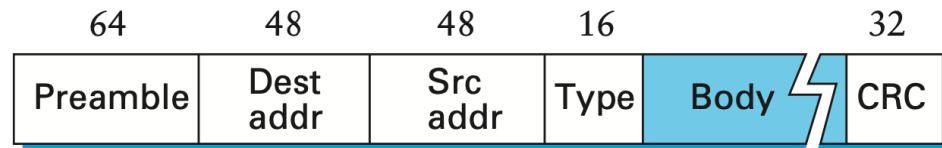


PROVIDE ABSTRACTION: ETHERNET (PHYSICAL LAYER)

- Ethernet Protocol (~80s)

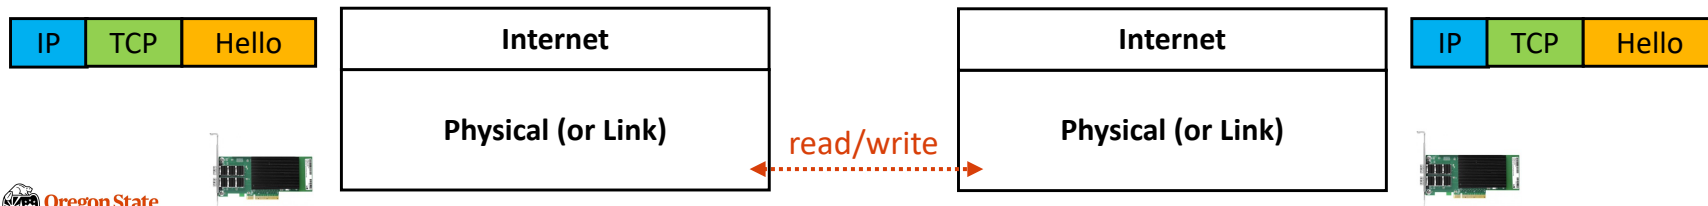
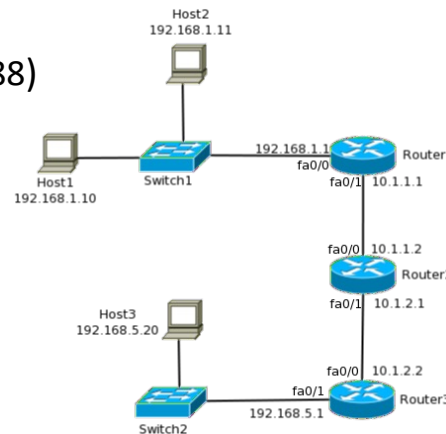
- Each network device (NIC) has 48-bit **MAC address**
- Each NIC is connected via Ethernet **cable**
- **ETH header** contains:

- (64 bit) Preamble (0x11111111... or a unique data)
- (48-bit) Destination MAC address
- (48-bit) Source MAC address
- (16-bit) Type
- (up to 1500 bytes) Data
- (32-bit) CRC for error correcting



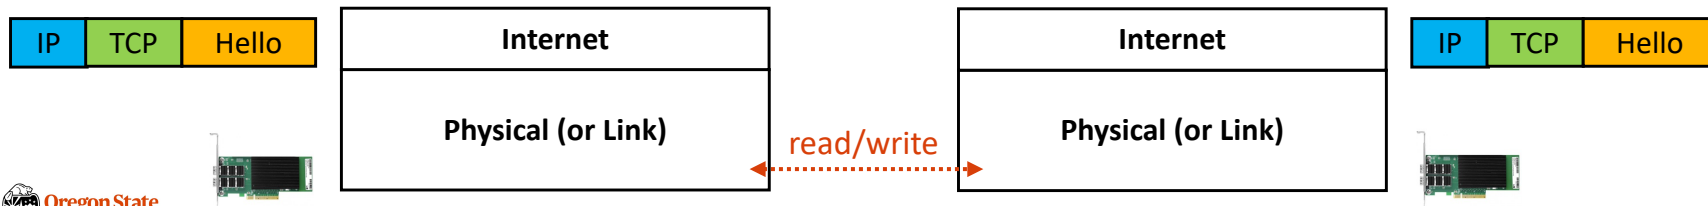
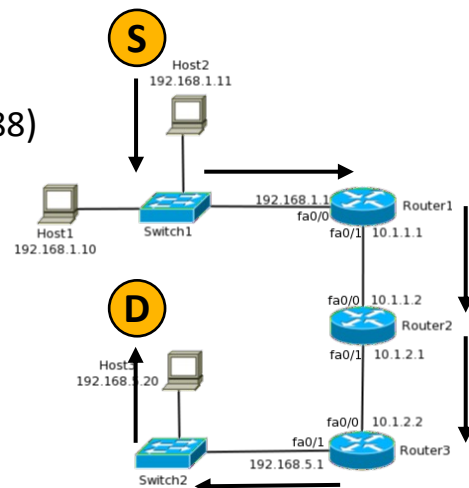
PROVIDE ABSTRACTION: IP LAYER

- Internet Protocol (IP)
 - IP allows us to connect multiple networks
 - Each host has a unique IP address
 - **IPv4**: 32-bit address (e.g., 147.56.28.101)
 - **IPv6**: 128-bit address (e.g., 2001:db8:3333:4444:5555:6666:7777:8888)



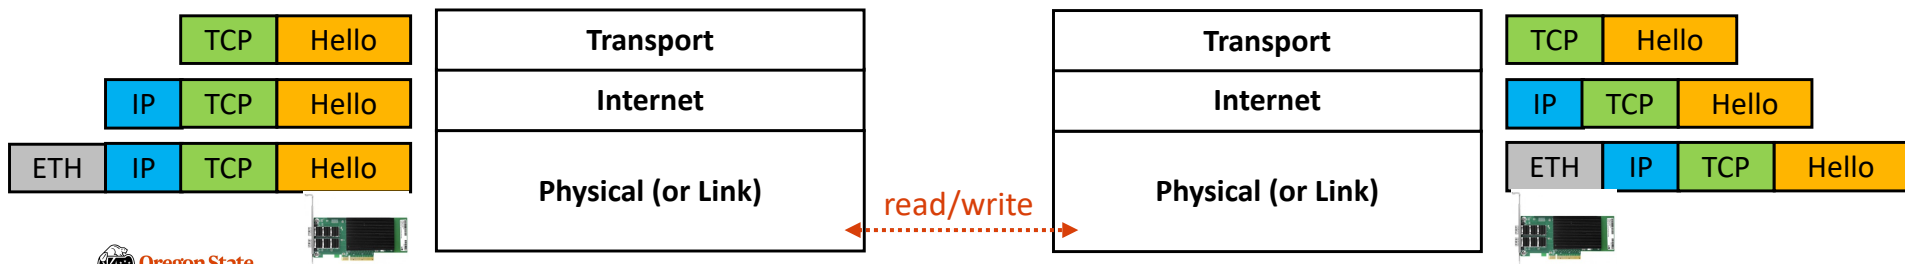
PROVIDE ABSTRACTION: IP LAYER

- Internet Protocol (IP)
 - IP allows us to connect multiple networks
 - Each host has a unique IP address
 - **IPv4**: 32-bit address (e.g., 147.56.28.101)
 - **IPv6**: 128-bit address (e.g., 2001:db8:3333:4444:5555:6666:7777:8888)
 - IP data (packets) is **routed** based on **destination IP**



PROVIDE ABSTRACTION: TRANSPORT LAYER

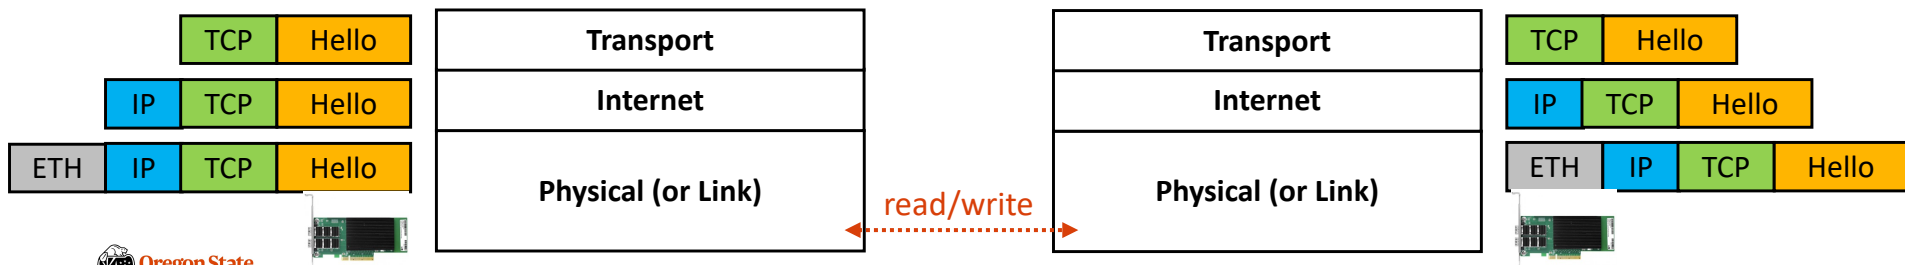
- TCP vs UDP **Protocol**
 - Transmission Control Protocol: **TCP** Packet
 - (16-bit, for each) Source and destination ports
 - (32-bit) Sequence number
 - (32-bit) Acknowledgement number
 - Others: flags, checksums, window-size, pointer, ...
 - User Datagram Protocol: **UDP** Packet
 - (16-bit, for each) Source and destination port
 - (16-bit, for each) Length and checksum



PROVIDE ABSTRACTION: TRANSPORT LAYER

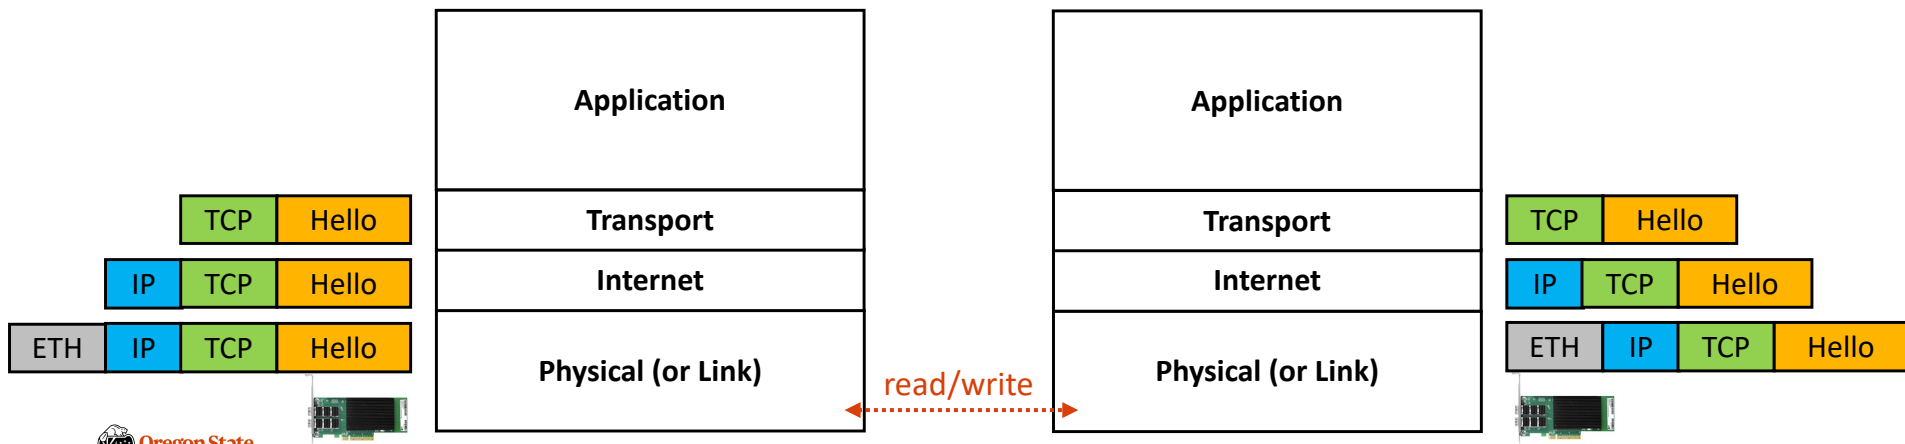
- TCP vs UDP **Protocol**

- TCP requires an established connection, but UDP is not (broadcast)
- TCP can use sequences, but UDP is not
- TCP is like a PIPE; data won't be lost, but UDP will (can lose data)
- TCP guarantees delivery, but UDP does not
- TCP is slower than UDP (suppose that we deliver all the packets)



PROVIDE ABSTRACTION: APPLICATION LAYER

- Application layer
 - Support various user-defined or OS-defined protocols (on top of TCP/UDP)
 - **TCP-based** : HTTPS, HTTP, SMTP, POP, FTP, ...
 - **UDP-based**: Video streaming, conferencing, DNS, VoIP, ...



TOPICS FOR TODAY

- Part III: Networking
 - Provide abstraction
 - OSI models
 - Packet encapsulation
 - Offer standard interface
 - RPC mechanisms (e.g., sockets)
 - Manage resources
 - Packet encapsulation in detail

PACKET ENCAPSULATION IN DETAIL

- Dive into the encapsulation
 - Hardware: NIC (and the network driver)
 - Physical: MAC address-based communication
 - Internet: IP address-based communication
 - Transport: Define protocols, *e.g.*, TCP or UDP
 - Application: Define custom protocols, *e.g.*, HTTP, HTTPS, FTP, etc.

PACKET ENCAPSULATION: NIC

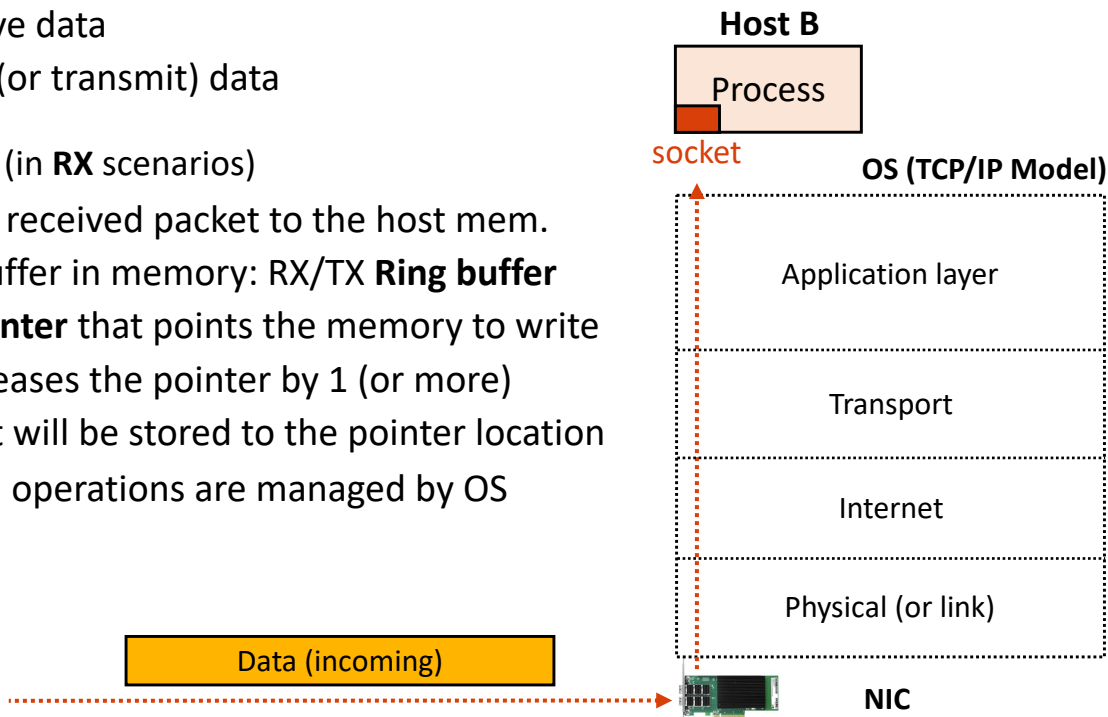
- Network Interface Card

- Networking terminology

- Receive (RX) : receive data
 - Transmit (TX): send (or transmit) data

- NIC and OS interaction (in **RX** scenarios)

- First, NIC copies the received packet to the host mem.
 - The OS has a buffer in memory: RX/TX **Ring buffer**
 - It also has **a pointer** that points the memory to write
 - Second, the OS increases the pointer by 1 (or more)
 - The next packet will be stored to the pointer location
 - The first and second operations are managed by OS



PACKET ENCAPSULATION: NIC

- Network Interface Card

- Networking terminology

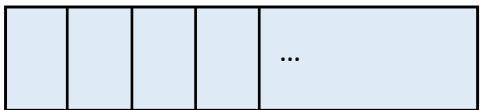
- Receive (RX) : receive data
 - Transmit (TX): send (or transmit) data

- NIC and OS interaction (in **RX** scenarios)

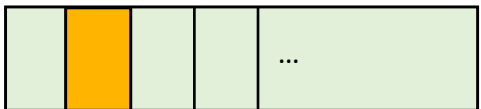
- First, NIC copies the received packet to the host mem.
 - The OS has a buffer in memory: RX/TX **Ring buffer**
 - It also has **a pointer** that points the memory to write
 - Second, the OS increases the pointer by 1 (or more)
 - The next packet will be stored to the pointer location
 - The first and second operations are managed by OS

OS (Device driver)

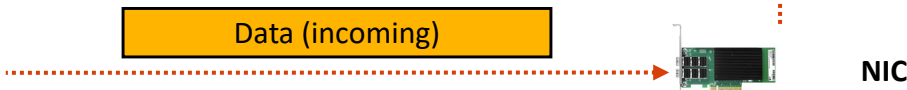
TX Ring Buffer



RX Ring Buffer



Pointer

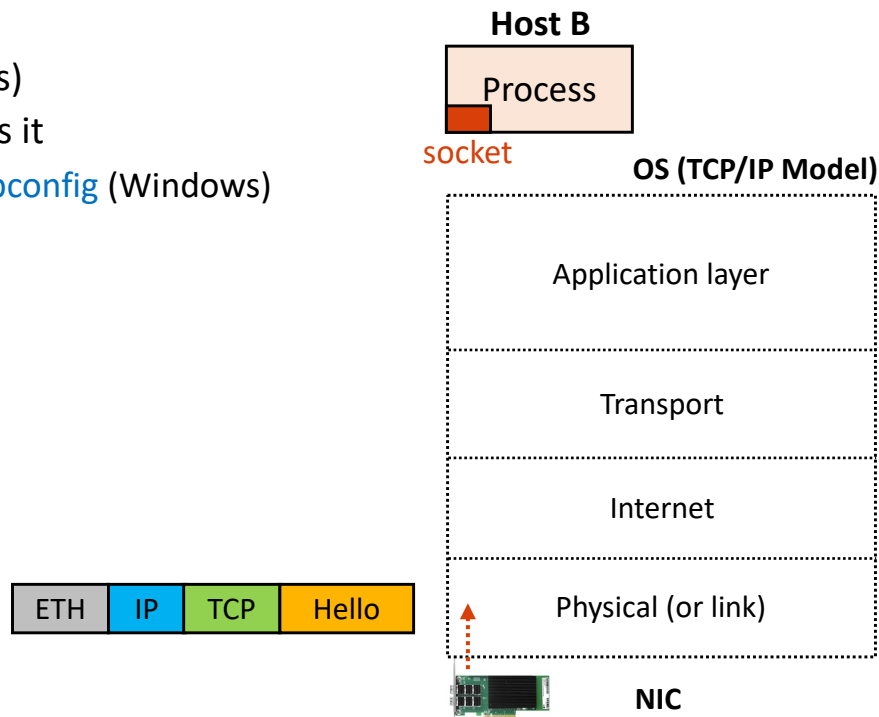


PACKET ENCAPSULATION IN DETAIL

- Dive into the encapsulation
 - Hardware: NIC (and the network driver)
 - Physical: MAC address-based communication
 - Internet: IP address-based communication
 - Transport: Define protocols, *e.g.*, TCP or UDP
 - Application: Define custom protocols, *e.g.*, HTTP, HTTPS, FTP, etc.

PACKET ENCAPSULATION: PHYSICAL (OR LINK) LAYER

- Physical (or link) layer
 - Ethernet Protocol (ETH)
 - Developed in 80s
 - Used for local area networks (LANs)
 - Use 48-bit MAC addresses; NIC has it
 - DIY check: `ifconfig` (Linux) and `ipconfig` (Windows)



PACKET ENCAPSULATION: PHYSICAL (OR LINK) LAYER

- Physical (or link) layer
 - Ethernet Protocol (ETH)
 - Developed in 80s
 - Used for local area networks (LANs)
 - Use 48-bit MAC addresses; NIC has it
 - DIY check: `ifconfig` (Linux) and `ipconfig` (Windows)

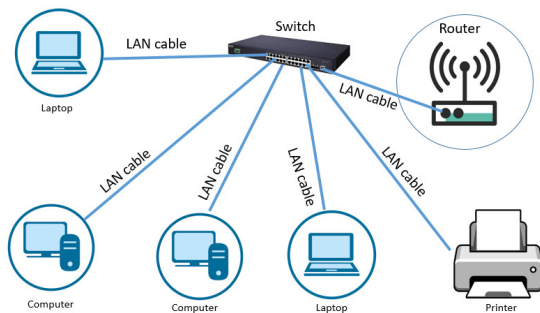
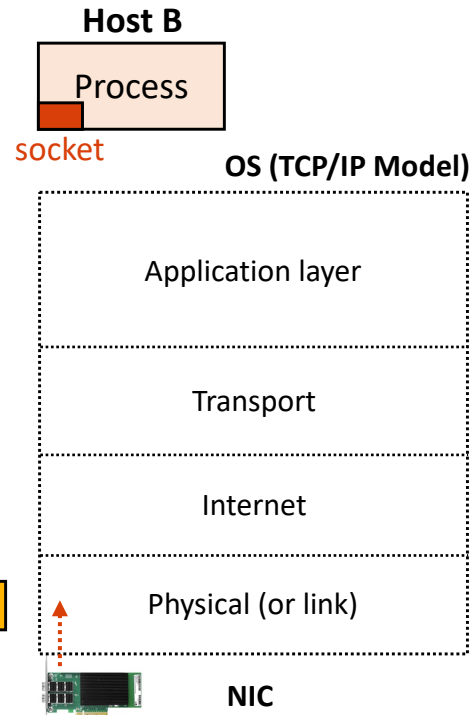
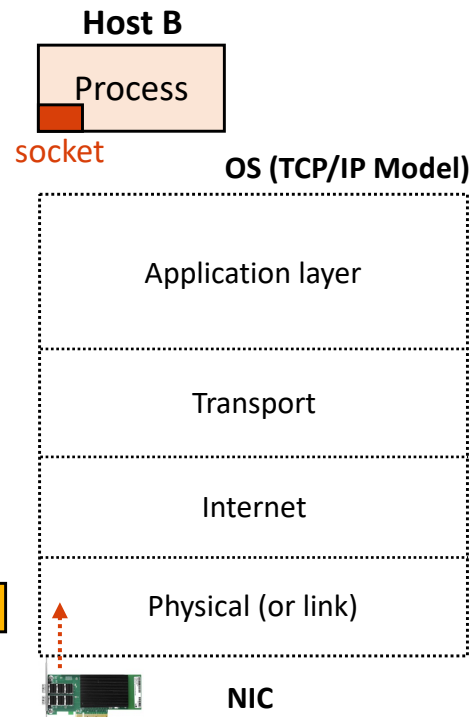


Illustration of a LAN
(Hosts are located quite nearby)



PACKET ENCAPSULATION: PHYSICAL (OR LINK) LAYER

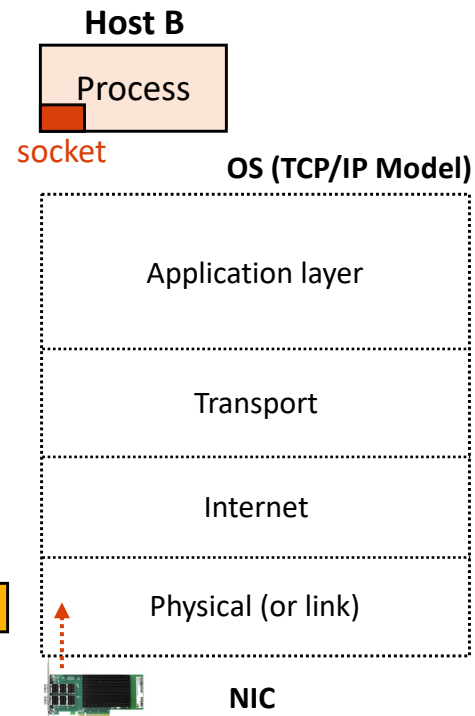
- Physical (or link) layer
 - Ethernet Protocol (ETH)
 - Developed in 80s
 - Used for local area networks (LANs)
 - Use 48-bit MAC addresses; NIC has it
 - DIY check: `ifconfig` (Linux) and `ipconfig` (Windows)
 - Ethernet packet contains:
 - Preamble (0x11111111... or a unique 64-bit data)
 - Destination MAC address (48-bit)
 - Source MAC address (48-bit)
 - Type (16-bit)
 - Data (variable ~1500 bytes)
 - CRC (32-bit; for error correcting)



PACKET ENCAPSULATION: PHYSICAL (OR LINK) LAYER

- Physical (or link) layer
 - Ethernet Protocol (ETH)
 - Developed in 80s
 - Used for local area networks (LANs)
 - Use 48-bit MAC addresses; NIC has it
 - DIY check: `ifconfig` (Linux) and `ipconfig` (Windows)
 - Ethernet packet contains:

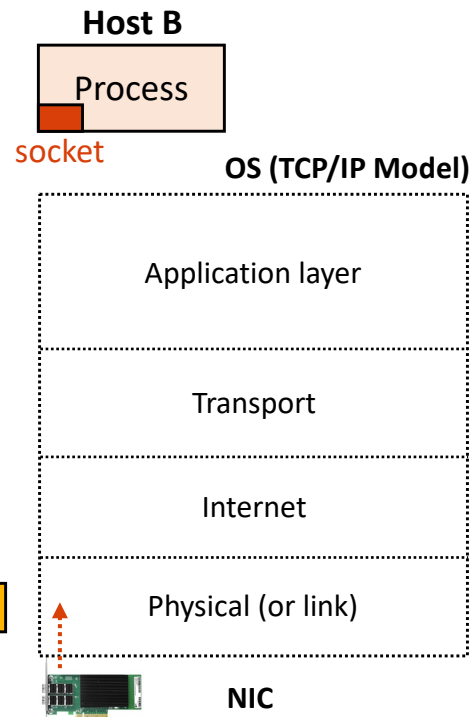
- NIC Manages →
- Preamble (0x11111111... or unique 64-bit data)
 - Destination MAC address (48-bit)
 - Source MAC address (48-bit)
 - Type (16-bit)
 - Data (variable ~1500 bytes)
 - CRC (32-bit; for error correcting)



PACKET ENCAPSULATION: PHYSICAL (OR LINK) LAYER

- Physical (or link) layer
 - Ethernet Protocol (ETH)
 - Developed in 80s
 - Used for local area networks (LANs)
 - Use 48-bit MAC addresses; NIC has it
 - DIY check: `ifconfig` (Linux) and `ipconfig` (Windows)
 - Ethernet packet contains:
 - Preamble (0x11111111... or unique 64-bit data)
 - Destination MAC address (48-bit)
 - Source MAC address (48-bit)
 - Type (16-bit)
 - Data (variable ~1500 bytes)
 - CRC (32-bit; for error correcting)

OS Manages



PACKET ENCAPSULATION: PHYSICAL (OR LINK) LAYER

- Physical (or link) layer
 - Ethernet Protocol (ETH)
 - Developed in 80s
 - Used for local area networks (LANs)
 - Use 48-bit MAC addresses; NIC has it
 - DIY check: `ifconfig` (Linux) and `ipconfig` (Win)
 - Ethernet packet contains:
 - Preamble (unique 64-bit data)
 - Destination MAC address (48-bit)
 - Source MAC address (48-bit)
 - Type (16-bit)
 - Data (variable ~1500 bytes)
 - CRC (32-bit; for error correcting)

```
/*  
 *      This is an Ethernet frame header.  
 */  
  
/* allow libcs like musl to deactivate this, glibc does not implement this. */  
#ifndef __UAPI_DEF_ETHHDR  
#define __UAPI_DEF_ETHHDR 1  
#endif  
  
#if __UAPI_DEF_ETHHDR  
struct ethhdr {  
    unsigned char    h_dest[ETH_ALEN];    /* destination eth addr */  
    unsigned char    h_source[ETH_ALEN];  /* source ether addr */  
    __be16           h_proto;              /* packet type ID field */  
} __attribute__((packed));  
#endif
```

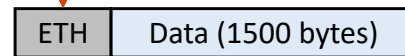
h_dest: Destination MAC address
h_source: Source MAC address
h_proto: Packet type ID field

ETH_ALEN: 6 bytes (48 bits)
(see the kernel code: [include/uapi/linux/if_ether.h](#))

PACKET ENCAPSULATION: PHYSICAL (OR LINK) LAYER

- Physical (or link) layer
 - Ethernet Protocol (ETH)
 - Developed in 80s
 - Used for local area networks (LANs)
 - Use 48-bit MAC addresses; NIC has it
 - DIY check: `ifconfig` (Linux) and `ipconfig` (Win)
 - Ethernet packet contains:
 - Preamble (unique 64-bit data)
 - Destination MAC address (48-bit)
 - Source MAC address (48-bit)
 - Type (16-bit)
 - Data (variable ~1500 bytes)
 - CRC (32-bit; for error correcting)

```
/*  
 *      This is an Ethernet frame header.  
 */  
  
/* allow libcs like musl to deactivate this, glibc does not implement this. */  
#ifndef __UAPI_DEF_ETHHDR  
#define __UAPI_DEF_ETHHDR 1  
#endif  
  
#if __UAPI_DEF_ETHHDR  
struct ethhdr {  
    unsigned char    h_dest[ETH_ALEN];    /* destination eth addr */  
    unsigned char    h_source[ETH_ALEN];  /* source ether addr */  
    __be16           h_proto;              /* packet type ID field */  
} __attribute__((packed));  
#endif
```



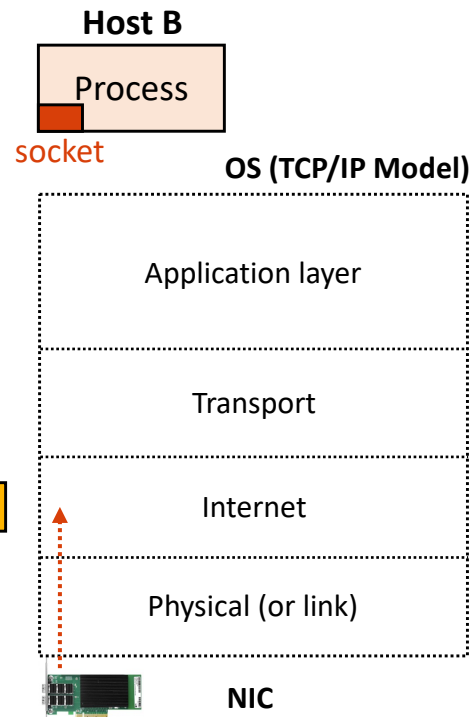
ETH_ETH_DATA_LEN: 1500
(see [include/uapi/linux/if_ether.h](#))

PACKET ENCAPSULATION IN DETAIL

- Dive into the encapsulation
 - Hardware: NIC (and the network driver)
 - Physical: MAC address-based communication
 - Internet: IP address-based communication
 - Transport: Define protocols, *e.g.*, TCP or UDP
 - Application: Define custom protocols, *e.g.*, HTTP, HTTPS, FTP, etc.

PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer
 - Internet Protocol (IP)
 - Use to connect multiple LANs
 - Use IP addresses to locate a host(s)
 - IPv4: 32-bit address, *e.g.*, 54.189.37.112
 - IPv6: 64-bit address, *e.g.*, 2001:0db8:85a3:0000:0000:8a2e:0370:7334



PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer
 - Internet Protocol (IP)
 - Use to connect multiple LANs
 - Use IP addresses to locate a host(s)
 - IPv4: 32-bit address, *e.g.*, 54.189.37.112
 - IPv6: 64-bit address, *e.g.*, 2001:0db8:85a3:0000:0000:8a2e:0370:7334

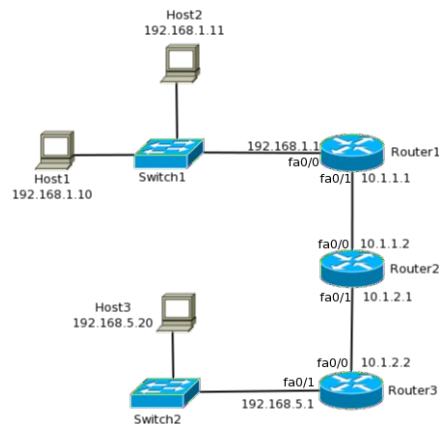
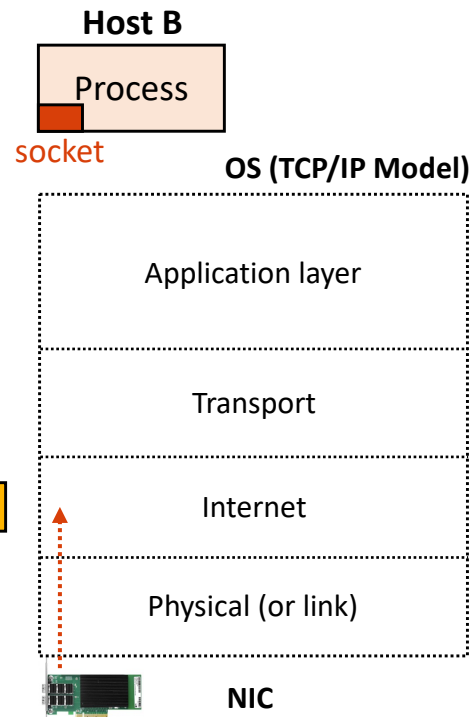
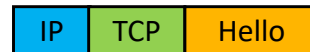


Illustration of the Internet
(Hosts are located remotely)

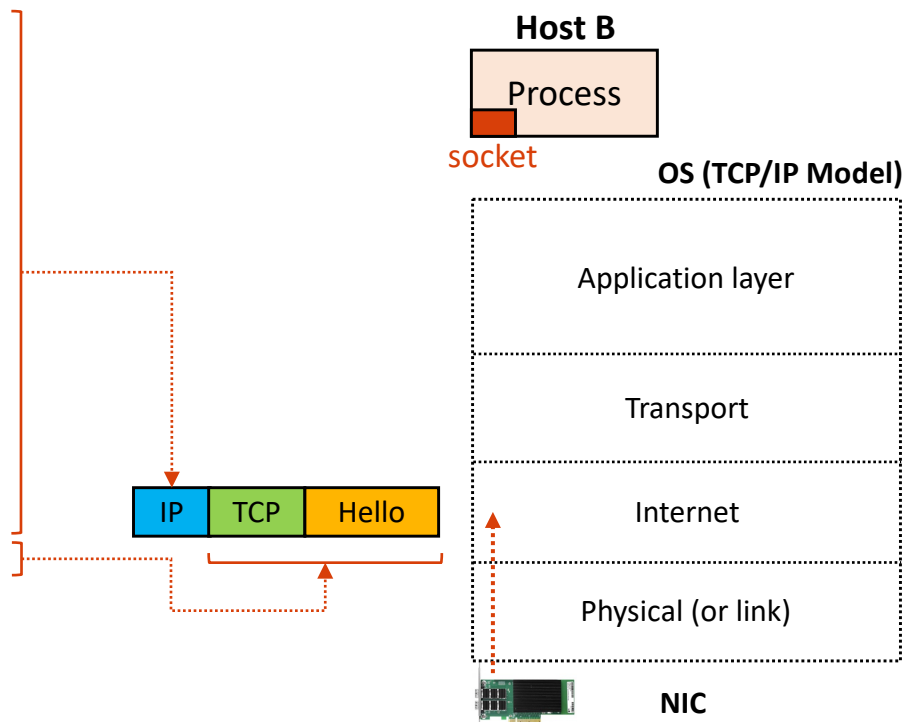


PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- IP packet contains:

- Version v4/v6 (8-bit)
 - Type of service (8-bit)
 - Total length (16-bit)
 - Identification (16-bit)
 - Frame offset (16-bit)
 - Time to live (8-bit)
 - Protocol (8-bit)
 - CRC (checksum) (16-bit)
 - Source IP address (32-bit)
 - Destination IP address (32-bit)
 - Data (1 ~ 65515 bytes)



PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer
 - IP packet **header** contains:
 - Version v4 (8-bit)
 - Type of service (8-bit)
 - Total length (16-bit)
 - Identification (16-bit)
 - Frame offset (16-bit)
 - Time to live (8-bit)
 - Protocol (8-bit)
 - CRC (checksum) (16-bit)
 - Source IP address (32-bit)
 - Destination IP address (32-bit)

IPv4 Header ([source code](#))

```
struct iphdr {  
#if defined(__LITTLE_ENDIAN_BITFIELD)  
    __u8    ihl:4,  
            version:4;  
#elif defined (__BIG_ENDIAN_BITFIELD)  
    __u8    version:4,  
            ihl:4;  
#else  
#error "Please fix <asm/byteorder.h>"  
#endif  
  
    __u8    tos;  
    __be16  tot_len;  
    __be16  id;  
    __be16  frag_off;  
    __u8    ttl;  
    __u8    protocol;  
    __sum16 check;  
    __be32  saddr;  
    __be32  daddr;  
    /*The options start here. */  
};
```

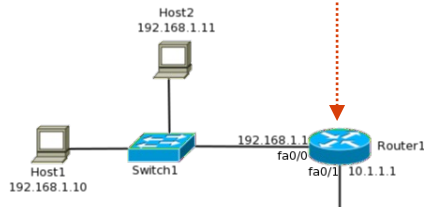
PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- IP packet header contains:

- Version v4 (8-bit)
 - Type of service (8-bit: [reference](#))
 - Total length (16-bit)
 - Identification (16-bit)
 - Frame offset (16-bit)
 - Time to live (8-bit)
 - Protocol (8-bit)
 - CRC (checksum) (16-bit)
 - Source IP address (32-bit)
 - Destination IP address (32-bit)

Priority of IP Packet
Important at the router
(CS 372 Topic)



IPv4 Header

```
struct iphdr {
#ifdef LITTLE_ENDIAN_BITFIELD
    __u8    ihl:4,
           version:4;
#elif defined (BIG_ENDIAN_BITFIELD)
    __u8    version:4,
           ihl:4;
#else
#error "Please fix <asm/byteorder.h>"
#endif
    __u8     tos;
    __be16   tot_len;
    __be16   id;
    __be16   frag_off;
    __u8     ttl;
    __u8     protocol;
    __sum16   check;
    __be32   saddr;
    __be32   daddr;
    /*The options start here. */
};
```

PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- IP packet header contains:

- Version v4 (8-bit)
 - Type of service (8-bit: [reference](#))
 - **Total length (16-bit)**→
 - Identification (16-bit)
 - Frame offset (16-bit)
 - Time to live (8-bit)
 - Protocol (8-bit)
 - CRC (checksum) (16-bit)
 - Source IP address (32-bit)
 - Destination IP address (32-bit)

Total packet length

Min : 21 bytes

Max: 65535 bytes

(20 bytes header +
1 ~ 65515 bytes data)

IPv4 Header

```
struct iphdr {
#ifdef __LITTLE_ENDIAN_BITFIELD
    __u8    ihl:4,
           version:4;
#elif defined (__BIG_ENDIAN_BITFIELD)
    __u8    version:4,
           ihl:4;
#else
#error "Please fix <asm/byteorder.h>"
#endif
    __u8     tos;
    __be16   tot_len;
    __be16   id;
    __be16   frag_off;
    __u8     ttl;
    __u8     protocol;
    __sum16   check;
    __be32   saddr;
    __be32   daddr;
    /*The options start here. */
};
```

PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- IP packet header contains:

- Version v4 (8-bit)
 - Type of service (8-bit: [reference](#))
 - Total length (16-bit)
 - Identification (16-bit)
 - Frame offset (16-bit)
 - Time to live (8-bit)
 - Protocol (8-bit)
 - CRC (checksum) (16-bit)
 - Source IP address (32-bit)
 - Destination IP address (32-bit)

Fragmentation

Oftentimes, we send multiple IP packets that sum up as a whole (ex. the entire data 4000 bytes, we send 4 1000-byte packets)

For example, 3rd IP packet's
ID: 3
Offset: 2000

IPv4 Header

```
struct iphdr {  
#if defined(__LITTLE_ENDIAN_BITFIELD)  
    __u8    ihl:4,  
            version:4;  
#elif defined (__BIG_ENDIAN_BITFIELD)  
    __u8    version:4,  
            ihl:4;  
#else  
#error "Please fix <asm/byteorder.h>"  
#endif  
    __u8    tos;  
    __be16  tot_len;  
    __be16  id;  
    __be16  frag_off;  
    __u8    ttl;  
    __u8    protocol;  
    __sum16 check;  
    __be32  saddr;  
    __be32  daddr;  
    /*The options start here. */  
};
```

PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- IP packet header contains:

- Version v4 (8-bit)
 - Type of service (8-bit: [reference](#))
 - Total length (16-bit)
 - Identification (16-bit)
 - Frame offset (16-bit)
 - Time to live (8-bit)
 - Protocol (8-bit)
 - CRC (checksum) (16-bit)
 - Source IP address (32-bit)
 - Destination IP address (32-bit)

Time-to-live (TTL)

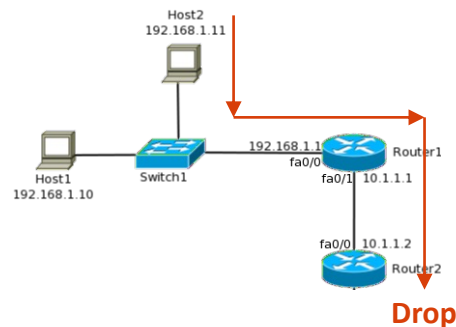
The number of routers a packet can maximally pass (# of “hops”)

Reason: drop packets that live too long in a network infrastructure

Example: on the right, setting TTL to 2 -> a packet drops at router 2

IPv4 Header

```
struct iphdr {  
#if defined(__LITTLE_ENDIAN_BITFIELD)  
    __u8    ihl:4,  
            version:4;  
#elif defined(__BIG_ENDIAN_BITFIELD)  
    __u8    version:4,  
            ihl:4;  
#else  
#error "Please fix <asm/byteorder.h>"  
#endif  
  
    __u8    tos;  
    __be16  tot_len;  
    __be16  id;  
    __be16  frag_off;  
    __u8    ttl;  
    __u8    protocol;  
    __sum16 check;  
    __be32  saddr;  
    __be32  daddr;  
    /*The options start here. */  
};
```



PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- IP packet header contains:

- Version v4 (8-bit)
 - Type of service (8-bit: [reference](#))
 - Total length (16-bit)
 - Identification (16-bit)
 - Frame offset (16-bit)
 - Time to live (8-bit)
 - **Protocol (8-bit)** →
 - CRC (checksum) (16-bit)
 - Source IP address (32-bit)
 - Destination IP address (32-bit)

IP Protocols

Examples:

TCP	6
UDP	7
ETHERNET	143

(see the list of IP protocols: [link](#))

IPv4 Header

```
struct iphdr {  
#if defined(__LITTLE_ENDIAN_BITFIELD)  
    __u8    ihl:4,  
            version:4;  
#elif defined (__BIG_ENDIAN_BITFIELD)  
    __u8    version:4,  
            ihl:4;  
#else  
#error "Please fix <asm/byteorder.h>"  
#endif  
  
    __u8    tos;  
    __be16  tot_len;  
    __be16  id;  
    __be16  frag_off;  
    __u8    ttl;  
    __u8    protocol;  
    __sum16 check;  
    __be32  saddr;  
    __be32  daddr;  
    /*The options start here. */  
};
```

PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- IP packet header contains:

- Version v4 (8-bit)
 - Type of service (8-bit: [reference](#))
 - Total length (16-bit)
 - Identification (16-bit)
 - Frame offset (16-bit)
 - Time to live (8-bit)
 - Protocol (8-bit)
 - CRC (checksum) (16-bit)→
 - Source IP address (32-bit)
 - Destination IP address (32-bit)

Checksum

The number to check the *integrity* of a packet while routing.

To make sure the packet is not manipulated, each router re-computes it and compares the checksum value with the stored one. (Q: what's the issue?)

IPv4 Header

```
struct iphdr {  
#if defined( __LITTLE_ENDIAN_BITFIELD )  
    __u8    ihl:4,  
            version:4;  
#elif defined( __BIG_ENDIAN_BITFIELD )  
    __u8    version:4,  
            ihl:4;  
#else  
#error "Please fix <asm/byteorder.h>"  
#endif  
  
    __u8    tos;  
    __be16  tot_len;  
    __be16  id;  
    __be16  frag_off;  
    __u8    ttl;  
    __u8    protocol;  
    __sum16  check;  
    __be32  saddr;  
    __be32  daddr;  
    /*The options start here. */  
};
```

PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- IP packet header contains:
 - Version v4 (8-bit)
 - Type of service (8-bit: [reference](#))
 - Total length (16-bit)
 - Identification (16-bit)
 - Frame offset (16-bit)
 - Time to live (8-bit)
 - Protocol (8-bit)
 - CRC (checksum) (16-bit)
 - Source IP address (32-bit) →
 - Destination IP address (32-bit)

IPv4 Header

```
struct iphdr {  
#if defined(__LITTLE_ENDIAN_BITFIELD)  
    __u8    ihl:4,  
            version:4;  
#elif defined (__BIG_ENDIAN_BITFIELD)  
    __u8    version:4,  
            ihl:4;  
#else  
#error "Please fix <asm/byteorder.h>"  
#endif  
  
    __u8    tos;  
    __be16  tot_len;  
    __be16  id;  
    __be16  frag_off;  
    __u8    ttl;  
    __u8    protocol;  
    __sum16 check;  
    __be32  saddr;  
    __be32  daddr;  
    /*The options start here. */  
};
```

Source / Destination IPs

Example:

Source : 53.59.125.98
Destination: 22.156.71.44

Router: each router has a routing table that stores the subnetwork addresses. It uses the subnetwork addr and source/destination IPs to decide where to send the packet

PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- Router mechanism:

- Router connects subnetworks
 - **Subnet:** a logical *division* of an IP net

Subnetwork examples:

Router B: 10.11.0.0/16 (65536 addr.)

Router C: 10.12.0.0/16 (same)

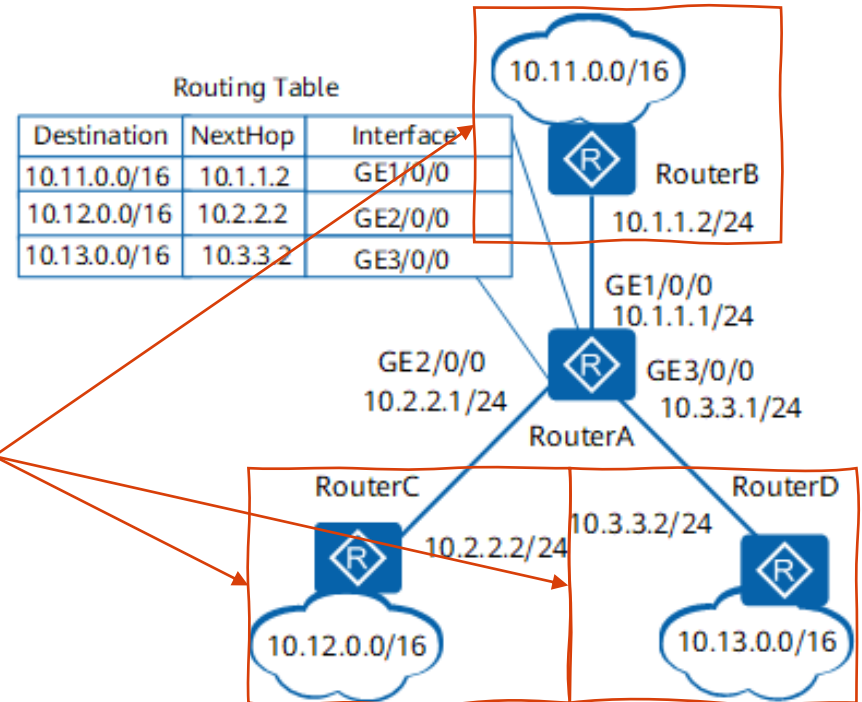
Router D: 10.13.0.0/16 (same)

CIDR Notations:

10.11.0.0 : base address

16 : # of leading bits we fix

10.11.0.0/16: 10.11.0.0 – 10.11.255.255



PACKET ENCAPSULATION: INTERNET LAYER

- Internet Layer

- Router mechanism:

- Router connects subnetworks
 - **Subnet:** a logical *division* of an IP net

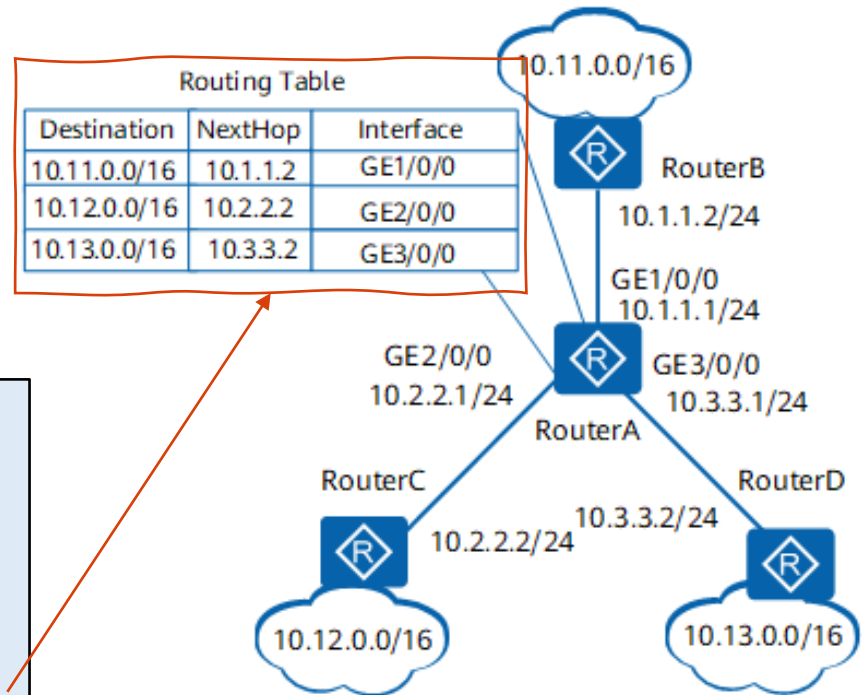
- Routing table:

- Describe the network destinations

In Router A: Suppose a packet comes from the source (10.12.1.45) moves to the destination (10.11.5.97)

Routing:

Host (src.) -> Router C
Router C -> Router A
Router A -> Router B (subnet: 10.11.0.0/16)
Router B -> Host (dest.)

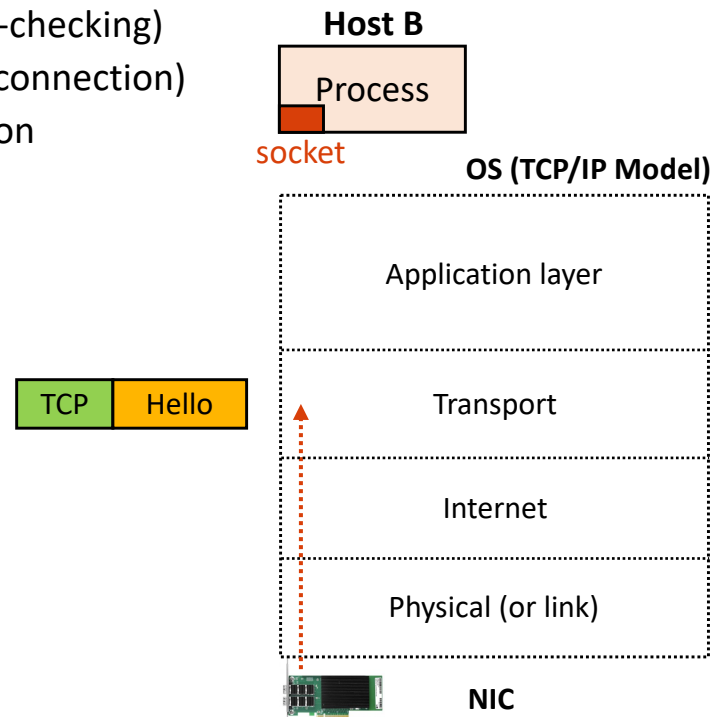


TOPICS FOR TODAY

- Dive into the encapsulation
 - Hardware: NIC (and the network driver)
 - Physical: MAC address-based communication
 - Internet: IP address-based communication
 - Transport: Define protocols, *e.g.*, TCP or UDP
 - Application: Define custom protocols, *e.g.*, HTTP, HTTPS, FTP, etc.

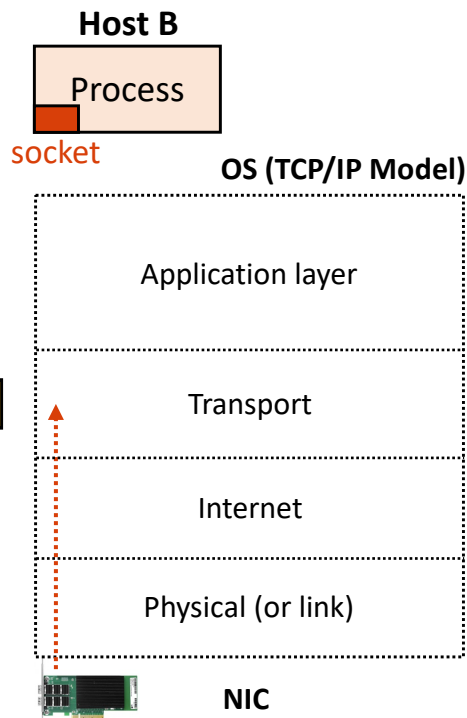
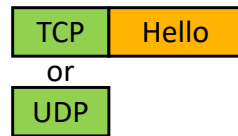
PACKET ENCAPSULATION: TRANSPORT LAYER

- Transport layer
 - Transmission Control Protocol (TCP)
 - **Reliable** communications (ordered and error-checking)
 - Connection oriented (first need to establish connection)
 - **3-way handshake** for establishing a connection



PACKET ENCAPSULATION: TRANSPORT LAYER

- Transport layer
 - Transmission Control Protocol (**TCP**)
 - Reliable communication (ordered and error-checking)
 - Connection oriented (first need to establish connection)
 - 3-way handshake for establishing a connection
 - User Datagram Protocol (**UDP**)
 - Simple **connectionless** communication (no handshake)
 - Less reliable communication (no order)
 - Useful for video/audio streaming [where losing packets is acceptable]



PACKET ENCAPSULATION: TRANSPORT LAYER

- Transport layer
 - Transmission Control Protocol (TCP)
 - Reliable communications (ordered and error-checking)
 - Connection oriented (first need to establish connection)
 - 3-way handshake for establishing a connection
 - TCP packet header contains:
 - Source ports (16-bit)
 - Destination ports (16-bits)
 - Sequence number (32-bit)
 - Acknowledgement number (32-bit)
 - Others (flags, checksums, window-size, pointer, ...)

TCP Header (source code)

```
struct tcphdr {
    __be16 source;
    __be16 dest;
    __be32 seq;
    __be32 ack_seq;
    #if defined(__LITTLE_ENDIAN_BITFIELD)
        __u16 res1:4,
            doff:4,
            fin:1,
            syn:1,
            rst:1,
            psh:1,
            ack:1,
            urg:1,
            ece:1,
            cwr:1;
    #elif defined(__BIG_ENDIAN_BITFIELD)
        __u16 doff:4,
            res1:4,
            cwr:1,
            ece:1,
            urg:1,
            ack:1,
            psh:1,
            rst:1,
            syn:1,
            fin:1;
    #else
        #error "Adjust your <asm/byteorder.h> defines"
    #endif
    __be16 window;
    __sum16 check;
    __be16 urg_ptr;
};
```

PACKET ENCAPSULATION: TRANSPORT LAYER

- Transport layer
 - Transmission Control Protocol (TCP)
 - Reliable communications (ordered and error-checking)
 - Connection oriented (first need to establish connection)
 - 3-way handshake for establishing a connection
 - TCP packet header contains:
 - Source ports (16-bit)
 - Destination ports (16-bits)
- **Source / Destination Ports**
 - IP addresses are in the IP packet header
 - TCP header only contains the port numbers (src/dest)

TCP Header (source code)

```
struct tcphdr {
    __be16 source;
    __be16 dest;
    __be32 seq;
    __be32 ack_seq;
#ifdef __LITTLE_ENDIAN_BITFIELD
    __u16 res1:4,
        doff:4,
        fin:1,
        syn:1,
        rst:1,
        psh:1,
        ack:1,
        urg:1,
        ece:1,
        cwr:1;
#elif defined(__BIG_ENDIAN_BITFIELD)
    __u16 doff:4,
        res1:4,
        cwr:1,
        ece:1,
        urg:1,
        ack:1,
        psh:1,
        rst:1,
        syn:1,
        fin:1;
#else
#error "Adjust your <asm/byteorder.h> defines"
#endif
    __be16 window;
    __sum16 check;
    __be16 urg_ptr;
};
```

PACKET ENCAPSULATION: TRANSPORT LAYER

- Transport layer
 - Transmission Control Protocol (TCP)
 - Reliable communications (ordered and error-checking)
 - Connection oriented (first need to establish connection)
 - 3-way handshake for establishing a connection
 - TCP packet header contains:
 - Source ports (16-bit)
 - Destination ports (16-bits)
 - Sequence number (32-bit)
 - Acknowledgement number (32-bit)
 - **Sequence number**
The **byte index** of the data a packet has (4103 / 11945)
 - **Ack number**
Receiver: # of packets remaining to receive
Sender : packet # to send next (to the receiver)

TCP Header (source code)

```
struct tcphdr {
    __be16 source;
    __be16 dest;
    __be32 seq;
    __be32 ack_seq;
    #if defined(__LITTLE_ENDIAN_BITFIELD)
        __u16 res1:4,
              doff:4,
              fin:1,
              syn:1,
              rst:1,
              psh:1,
              ack:1,
              urg:1,
              ece:1,
              cwr:1;
    #elif defined(__BIG_ENDIAN_BITFIELD)
        __u16 doff:4,
              res1:4,
              cwr:1,
              ece:1,
              urg:1,
              ack:1,
              psh:1,
              rst:1,
              syn:1,
              fin:1;
    #else
        #error "Adjust your <asm/byteorder.h> defines"
    #endif
    __be16 window;
    __sum16 check;
    __be16 urg_ptr;
};
```


PACKET ENCAPSULATION: TRANSPORT LAYER

- Transport layer
 - Transmission Control Protocol (TCP)
 - Reliable communications (ordered and error-checking)
 - Connection oriented (first need to establish connection)
 - 3-way handshake for establishing a connection
 - TCP packet header contains:
 - Source ports (16-bit)
 - Destination ports (16-bits)
 - Sequence number (32-bit)
 - Acknowledgement number (32-bit)
 - Others (flags, checksums, window-size, pointer, ...)

Other fields

Refer to this Wikipedia article ([link](#))

TCP Header ([source code](#))

```
struct tcphdr {
    __be16 source;
    __be16 dest;
    __be32 seq;
    __be32 ack_seq;
#ifdef __LITTLE_ENDIAN_BITFIELD
    __u16 res1:4,
        doff:4,
        fin:1,
        syn:1,
        rst:1,
        psh:1,
        ack:1,
        urg:1,
        ece:1,
        cwr:1;
#elif defined(__BIG_ENDIAN_BITFIELD)
    __u16 doff:4,
        res1:4,
        cwr:1,
        ece:1,
        urg:1,
        ack:1,
        psh:1,
        rst:1,
        syn:1,
        fin:1;
#else
#error "Adjust your <asm/byteorder.h> defines"
#endif
    __be16 window;
    __sum16 check;
    __be16 urg_ptr;
};
```

PACKET ENCAPSULATION: TRANSPORT LAYER

- Transport layer
 - User Datagram Protocol (UDP)
 - Simple connectionless communication (no handshake)
 - Less reliable communication (no order)
 - Useful for video/audio streaming

– UDP packet header contains:

- Source port (16-bit)
- Destination port (16-bits)

- **Source / Destination Ports**

IP addresses are in the IP packet header
UDP header only contains the port numbers like TCP

```
struct udphdr {  
    __be16 source;  
    __be16 dest;  
    __be16 len;  
    __sum16 check;  
};
```

PACKET ENCAPSULATION: TRANSPORT LAYER

- Transport layer
 - User Datagram Protocol (UDP)
 - Simple connectionless communication (no handshake)
 - Less reliable communication (no order)
 - Useful for video/audio streaming
 - UDP packet header contains:
 - Source port (16-bit)
 - Destination port (16-bits)
 - Packet length (16-bit)
 - Checksum (16-bits)

Length

Total UDP packet size (max. 65515 bytes)

Note: 65535 – 8-byte UDP header – 20-byte IP header

Checksum

Optional field as IP header already contains this data

```
struct udphdr {  
    __be16 source;  
    __be16 dest;  
    __be16 len;  
    __sum16 check;  
};
```

TOPICS FOR TODAY

- Dive into the encapsulation
 - Hardware: NIC (and the network driver)
 - Physical: MAC address-based communication
 - Internet: IP address-based communication
 - Transport: Define protocols, *e.g.*, TCP or UDP
 - Application: Define custom protocols, *e.g.*, HTTP, HTTPS, FTP, etc.
(**CS 372 Topic – Have more fun in this class**)

TOPICS FOR TODAY

- Part III: Networking
 - Provide abstraction
 - OSI models
 - Packet encapsulation
 - Offer standard interface
 - RPC mechanisms (e.g., sockets)
 - Manage resources
 - Packet encapsulation in detail

Thank You!

M/W 12:00 – 1:50 PM (LINC #200)

Sanghyun Hong

sanghyun.hong@oregonstate.edu



Oregon State
University

SAIL

Secure AI Systems Lab